

Parallelism in Software

Minsoo Ryu

Department of Computer Science and Engineering
Hanyang University



Outline

- 1 Parallelism in Software
- 2 Creating a Multicore Program
- 3 Multicore Design Patterns
- 4 Q & A

Types of Parallelism

➤ Parallelism in hardware (implicit parallelism)

- **Pipelining**
 - **Superscalar, VLIW**
 - **SIMD processing**
 - **HW multithreading**
- } instruction level parallelism (ILP)
- } data level parallelism (DLP)
- } thread level parallelism (TLP)

➤ Parallelism in software

- **Data parallelism**
- **Task parallelism**
 - Also known as control parallelism or function parallelism


Parallelism in Software: Data Parallelism

➤ Data parallelism

- A data set can be partitioned into multiple subsets
- A set of tasks perform the same computation, but operate on different data

➤ Consider applying a function square to the elements of an array A and then computing its sum

A = array of all data
Asqr = map(square, A)
s = sum(Asqr)



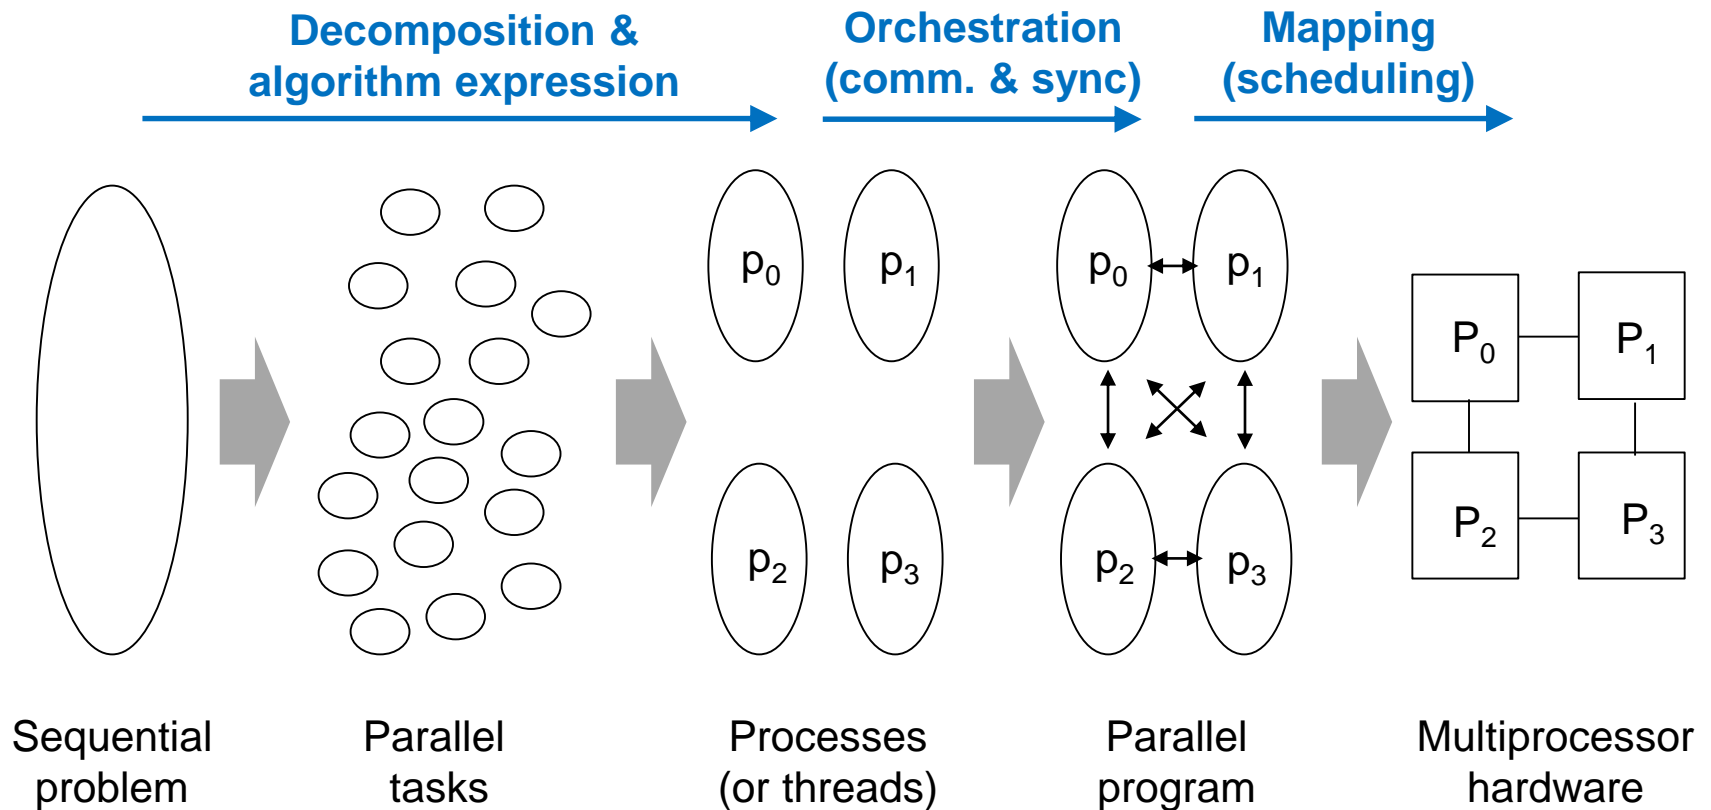
Parallelism in Software: Task Parallelism

➤ Task parallelism

- An application can be partitioned into multiple different tasks that can be executed simultaneously
- Task parallelism often involves data parallelism



Creating a Multicore Program



Four Design Spaces

Algorithmic Expression

- **Finding concurrency**
 - Decompose the problem (expose concurrent tasks)

- **Algorithm structure**
 - Map tasks to processes to exploit parallel architecture

Software Construction

- **Supporting structures**
 - Code and data structuring


- **Implementation mechanisms**
 - Low level mechanisms used to write parallel programs

Problem Decomposition

- Decompose the problem into smaller tasks
 - Determine which can be done in parallel with each other, and which must be done in some order

- Conceptualize a decomposition as a task dependency graph with:

- nodes corresponding to tasks
- edges indicating dependencies



- A given problem may be decomposed into tasks in many different ways
 - Tasks may be of same, different, or indeterminate sizes

Multicore Design Patterns

➤ What is a design pattern?

- A design pattern is a general reusable solution to a commonly occurring problem within a given context
- A design pattern is not a finished solution, but a description or template (a formalized best practice)

➤ Books on design patterns

- “Design Patterns: Elements of Reusable Object-Oriented Software,” E. Gamma, R. Helm, R. Johnson and J. Vlissides
- “Patterns for Parallel Programming,” T.G. Mattson, B.A. Sanders, and B.L. Massingill

Input/Output Data Decomposition

- Data decomposition for matrix multiplication
 - Matrices are partitioned into four submatrices

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

(a)

$$C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$$

$$C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$$

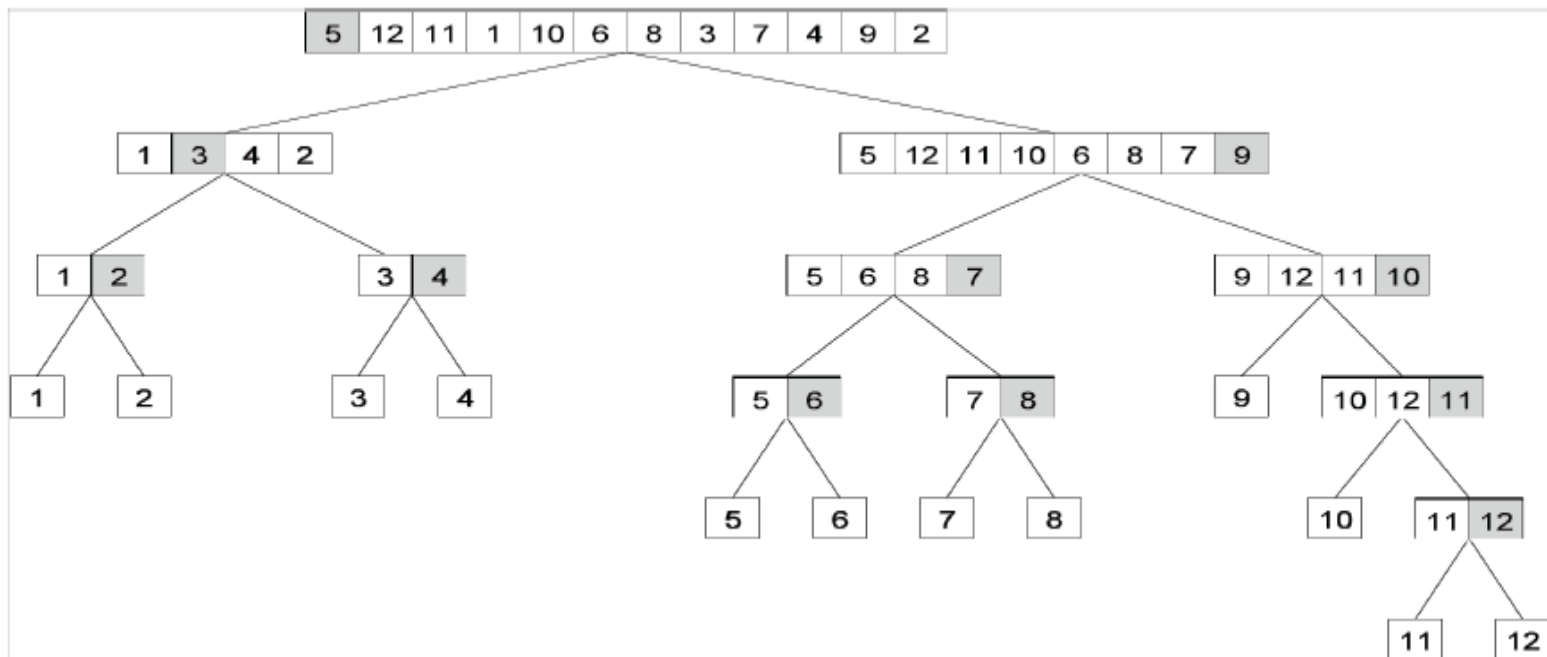
$$C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$$

$$C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$$

(b)


Recursive Decomposition

- Recursive decomposition for Quicksort
 - dived-and-conquer algorithm



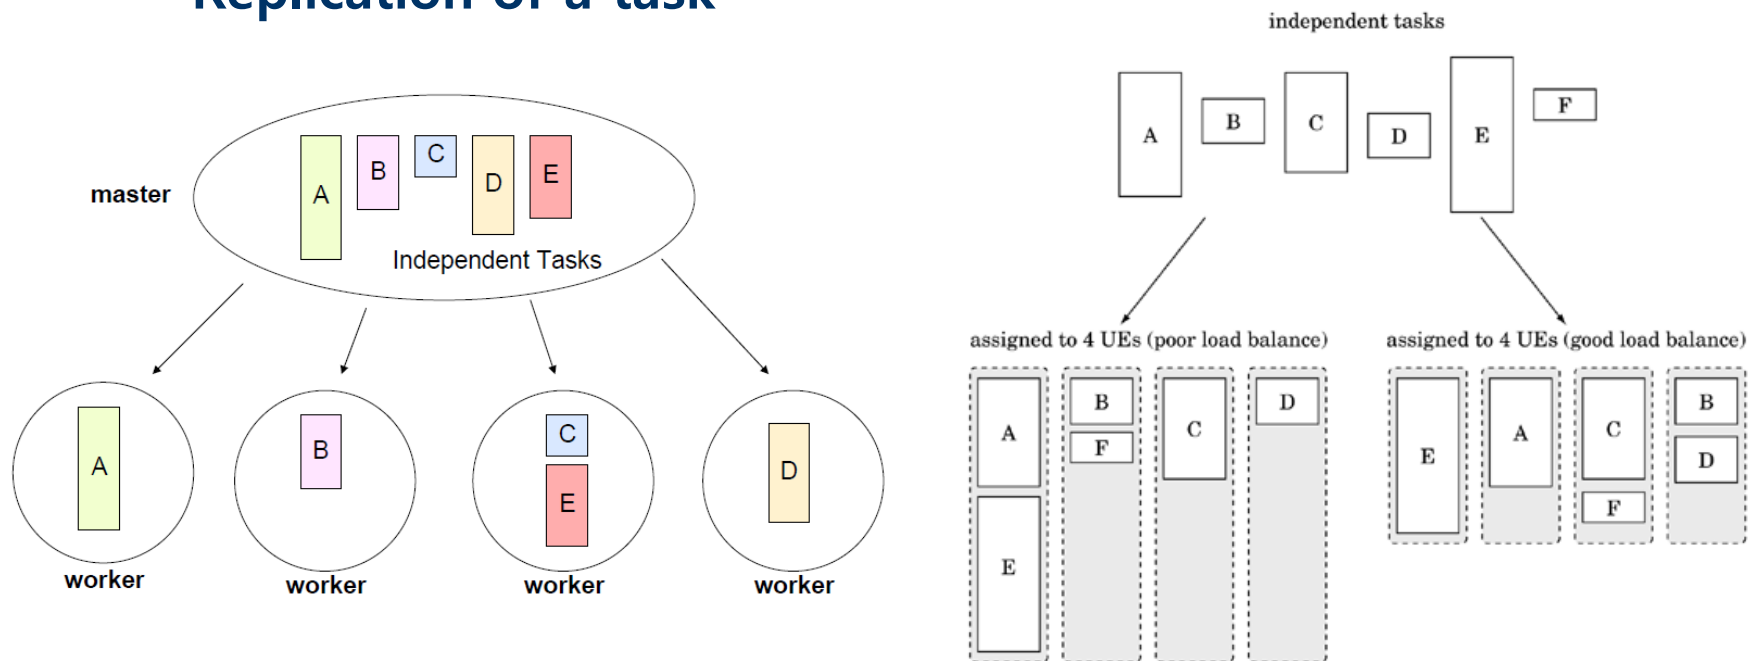
Hybrid Decomposition

- Hybrid decomposition for finding the minimum
 - A mix of decomposition techniques is often needed



Direct Task Decomposition


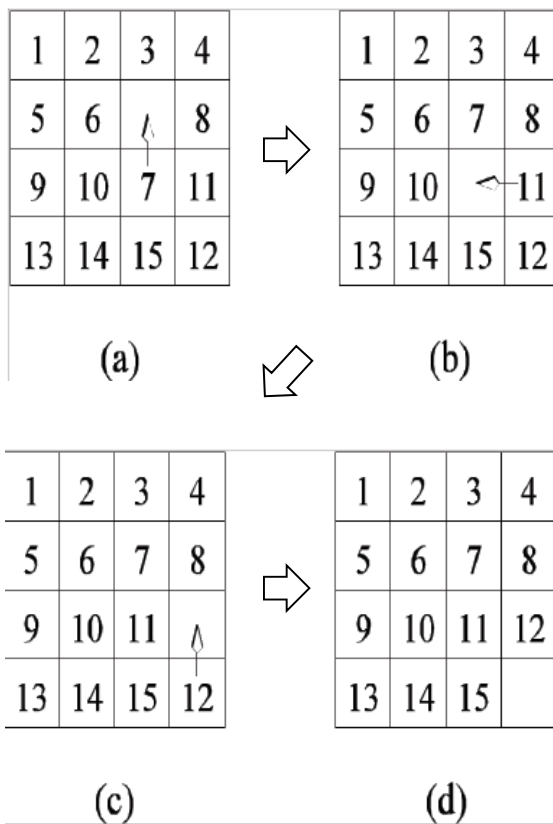
- Organize the system with independent tasks
 - Decomposition into independent tasks
 - Replication of a task



- Tasks can be executed in a master/worker manner

Exploratory Decomposition

- Exploratory decomposition for a tile puzzle
 - Search of a state space of solutions




Loop Parallelism

➤ Loop parallelism pattern

- Many programs are expressed using iterative constructs
- We can divide iterations into multiple sets of iterations and execute them in parallel

```
for (i = 0; i < 12; i++)  
    C[i] = A[i] + B[i];
```




Fork and Join

➤ Fork/Join pattern

- Parent task creates new tasks (fork) then waits until they complete (join) before continuing on with the computation


```
for (i = 0; i < 12; i++)  
    C[i] = A[i] + B[i];
```




Pipeline

➤ Pipeline pattern

- Tasks are organized in terms of the flow of data
- Examples
 - Ray tracing, molecular dynamics simulation, video codec



Linear pipeline

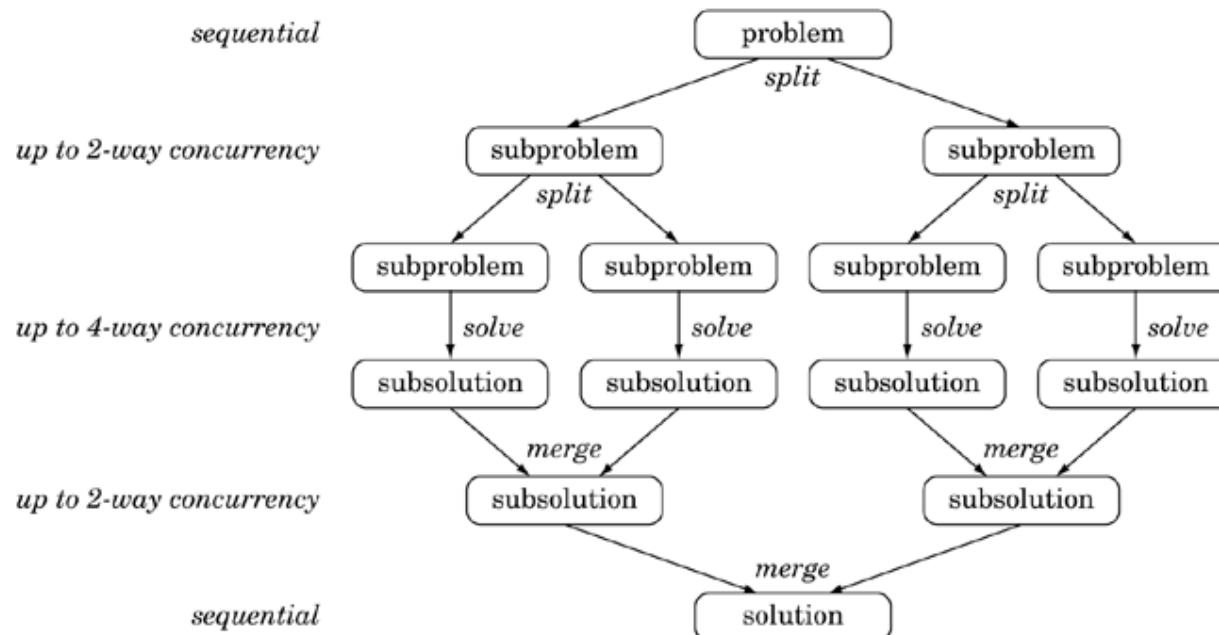


Nonlinear pipeline

Divide and Conquer

➤ Divide and conquer pattern

- A problem is solved by splitting it into a number of smaller subproblems, solving them independently, and merging the subsolutions
- Examples: Merge sort, matrix diagonalization



Example 1. Database Query Processing

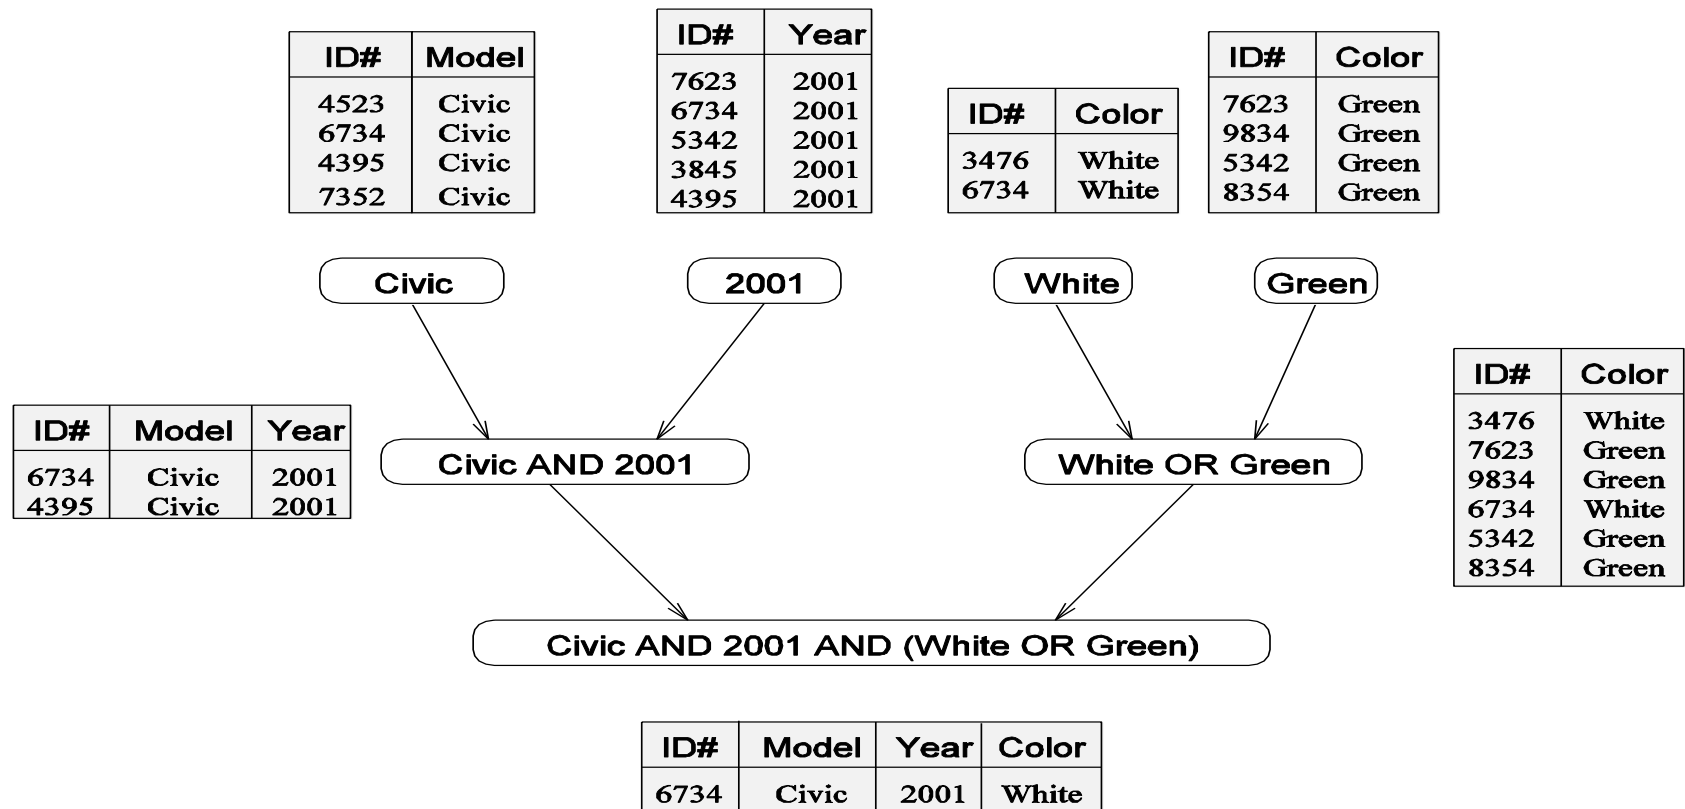
➤ Consider the following query

- **MODEL = "CIVIC" AND YEAR = 2001 AND COLOR = "GREEN" OR COLOR = "WHITE "**

ID#	Model	Year	Color	Dealer	Price
4523	Civic	2002	Blue	MN	\$18,000
3476	Corolla	1999	White	IL	\$15,000
7623	Camry	2001	Green	NY	\$21,000
9834	Prius	2001	Green	CA	\$18,000
6734	Civic	2001	White	OR	\$17,000
5342	Altima	2001	Green	FL	\$19,000
3845	Maxima	2001	Blue	NY	\$22,000
8354	Accord	2000	Green	VT	\$18,000
4395	Civic	2001	Red	CA	\$17,000
7352	Civic	2002	Red	WA	\$18,000

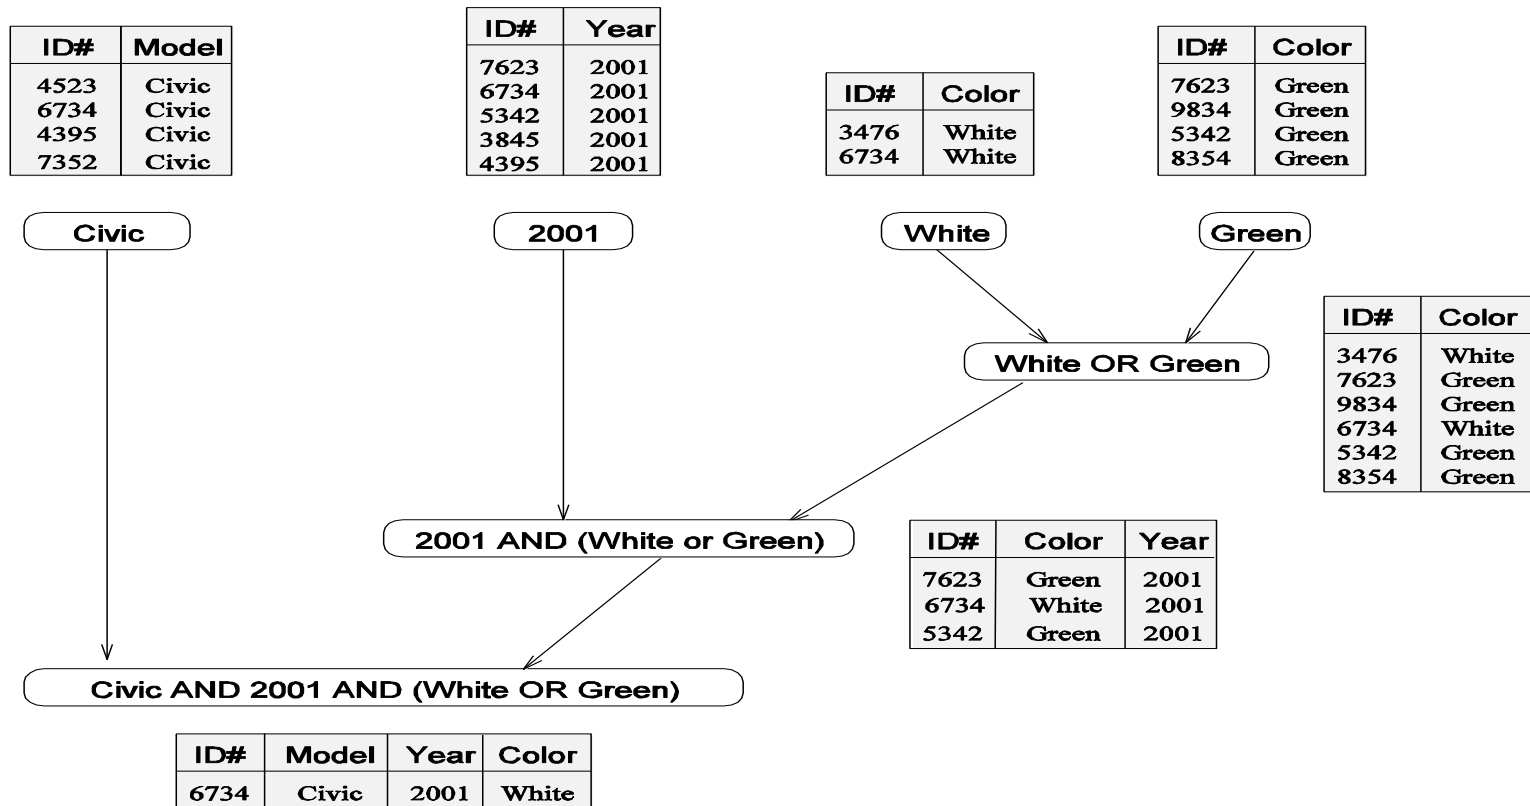
Example 1. Database Query Processing

➤ One possible decomposition




Example 1. Database Query Processing

➤ A different approach



Example 2. Monte Carlo Pi Calculation

- Estimate Pi by throwing darts at a unit square
- Calculate percentage that fall in the unit circle
 - Area of square = $r^2 = 1$
 - Area of circle quadrant = $\frac{1}{4} * \pi r^2 = \pi/4$
- Randomly throw darts at x, y positions
- If $x^2 + y^2 < 1$, then point is inside circle
- Compute ratio:
 - # points inside / # points total
 - $\pi = 4 * \text{ratio}$



Example 2. Monte Carlo Pi Calculation

```
main(int argc, char **argv) {  
    int i, hits, trials = 0;  
    double pi;  
  
    if (argc != 2) trials = 1000000;  
    else trials = atoi(argv[1]);  
  
    srand(0); // see hw tutorial  
  
    for (i=0; i < trials; i++) hits += hit();  
    pi = 4.0*hits/trials;  
    printf("PI estimated to %f.", pi);  
}
```

Example 2. Monte Carlo Pi Calculation

```
int hit() {  
    int const rand_max = 0xFFFFFFFF;  
    double x = ((double) rand()) / RAND_MAX;  
    double y = ((double) rand()) / RAND_MAX;  
    if ((x*x + y*y) <= 1.0) {  
        return(1);  
    } else {  
        return(0);  
    }  
}
```