

Synchronization for Concurrent Tasks

Minsoo Ryu

Department of Computer Science and Engineering
Hanyang University



Outline

1 Race Condition and Critical Section

Page X

2 Algorithmic Approaches

Page X

3 Hardware Support

Page X

4 OS Primitives

Page X

5 Q & A

Page X

Race Condition

- The situation where several processes access – and manipulate shared data concurrently
 - The final value of the shared data depends upon which process finishes last

Thread A

```
item nextProduced;  
  
if (user_wants_to_write == 1) {  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Thread B

```
item nextConsumed;  
  
If (user_wants_to_read == 1) {  
    while (counter == 0)  
        ; /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```

General Solution Structure

- General structure of process P_i (other process P_j)

do {

entry section

critical section

exit section

remainder section

} while (1);

- Processes may share some common variables to synchronize their actions

Example

Thread A

```
item nextProduced;

if (user_wants_to_write == 1) {
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    pthread_mutex_lock(&mutex);
    counter++;
    pthread_mutex_lock(&mutex);
}
```

Thread B

```
item nextConsumed;

If (user_wants_to_read == 1) {
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    pthread_mutex_lock(&mutex);
    counter--;
    pthread_mutex_lock(&mutex);
}
```

Three Approaches to Critical Sections

- **Algorithmic approaches**
 - Algorithmically solves the critical section problem without using any special HW and OS support
- **Hardware support**
 - Use special hardware support to achieve atomicity
 - Interrupt disabling, Test and Set instruction, Swap instruction
- **OS primitives**
 - Use OS primitives
 - Semaphore, mutex, spin lock, reader-writer lock, ...

Algorithmic Approach: Bakery Algorithm by Lamport

- Before entering its critical section, process receives a number
 - Holder of the smallest number enters the critical section
- If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first
- The numbering scheme always generates numbers in increasing order of enumeration; i.e.,
1,2,3,3,3,3,4,5...

Bakery Algorithm by Lamport

➤ Process P_i

- Initially, $\text{number}[i] = 0$

```

do {
    choosing[i] = true;
    number[i] = max (number[0], number[1] , ..., number [n - 1]) + 1;
    choosing[i] = false;
    for (j = 0; j < n; j++) {
        while (choosing[ j ]) ;
        while ((number[ j ] != 0) && ((number[ j ] < (number[i]) || (j < i)) ;
    }
    critical section
    number[i] = 0;
    remainder section
} while (1);
  
```

ties are broken by comparing j and i

Hardware Approach: Synchronization by Interrupt Disabling

- **Guarantee atomicity for accessing a critical section**
 - Forbid interrupts to occur in a critical section
 - The current sequence of instructions would be allowed to execute in order without preemption (atomicity)
 - No unexpected modifications could be made

- **Unfortunately, this solution is not feasible in a multiprocessor environment**
 - Disabling interrupts on a multiprocessor requires significant time to propagate the message to all the processors

Synchronization by Test-and-Set

- Many machines provide special HW instructions
 - To allow us to test and modify the content of a word, or
 - To swap the contents of two words, atomically
- Test and set instruction

```
boolean TestAndSet(boolean &target) {  
    boolean rv = target;  
    target = true;  
  
    return rv;  
}
```

Mutual Exclusion with Test-and-Set

- **Shared data:**
boolean lock = false;

- **Process P_i**
do {
 while (TestAndSet(lock)) ;
 critical section
 lock = false;
 remainder section
} while(1)

OS Primitives for Critical Sections

- OS provide several primitives for critical sections
 - Mutual exclusion, progress, and bounded waiting are guaranteed by OS
 - Most popular primitives are semaphores and mutexes

```
int main()
{
    sem_t semaphore;
    int count = 0;

    sem_init( &semaphore, 0, 1 );
    sem_wait( &semaphore );
    count++;
    sem_post( &semaphore );
    sem_destroy( &semaphore );
}
```

```
void * count( void * param )
{
    for ( int i=0; i<100; i++ )
    {
        pthread_mutex_lock( &mutex );
        counter++;
        printf( "Count = %i\n", counter );
        pthread_mutex_unlock( &mutex );
    }
}
```

Semaphores vs. Mutexes

➤ Semaphore

- A semaphore is a counter that can be either incremented or decremented
- Two types of semaphores
 - Counting semaphores (integer numbers)
 - Binary semaphores (0 or 1)
- Semaphores can be shared by different processes

➤ Mutex

- A binary semaphore
- Mutex has a sense of ownership
 - Only the task that has lock the mutex can unlock it
- By default, a mutex can be shared between threads

Deadlock

- Semaphores or mutexes must be used carefully
 - Otherwise, a deadlock or a livelock can happen

Thread 1

```
void update1()
{
    acquire(A);
    acquire(B); <<< Thread 1
                  waits here

    variable1++;
    release(B);
    release(A);
}
```

Thread 2

```
void update2()
{
    acquire(B);
    acquire(A); <<< Thread 2
                  waits here

    variable1++;
    release(B);
    release(A);
}
```

Livelock

- Threads can be trapped in a livelock of constantly acquiring and releasing mutexes

```
int done=0;
while (!done)
{
    acquire(A);
    if ( canAcquire(B) )
    {
        variable1++;
        release(B);
        release(A);
        done=1;
    }
    else
    {
        release(A);
    }
}
```

```
int done=0;
while (!done)
{
    acquire(B);
    if ( canAcquire(A) )
    {
        variable2++;
        release(A);
        release(B);
        done=1;
    }
    else
    {
        release(B);
    }
}
```

Spin Locks

- **Spin locks are essentially mutex locks**
 - Tasks waiting for mutex locks can sleep or spin
 - Tasks using a spin lock keeps trying to acquire the lock without sleeping

- **Advantage of spin locks**
 - Tasks will acquire the lock as soon as it is released
 - For a mutex lock that sleeps, the task needs to be woken by the operating system before it can get the lock

- **Disadvantage of spin locks**
 - A spin lock will be monopolizing the CPU

POSIX Spinlocks

```
pthread_spinlock_t lock;

void lockandunlock()
{
    int i = 10000;
    while ( i>0 )
    {
        pthread_spin_lock( &lock );
        i--;
        pthread_spin_unlock( &lock );
    }
}

int main()
{
    pthread_spin_init( &lock, PTHREAD_PROCESS_PRIVATE );
    lockandunlock();
    pthread_spin_destroy( &lock );
}
```

Synchronization in Real World

- OS primitives often serve as a powerful tool for many applications
- However, there are many real world problems that are hard to solve only with OS primitives
 - We often need a more complex approach
 - One difficult problem is the “Readers-Writers Synchronization”

The First Readers-Writers Problem

- The first readers-writers problem
 - Give preferential treatment to readers
 - Writers may suffer unbounded waiting

- Shared data
 - semaphore S , wrt;
 - **Initially**, $S = 1$, wrt = 1, readcount = 0


The First Readers-Writers Problem

```
/* Writer */  
wait(wrt);  
...  
writing is performed  
...  
signal(wrt);
```

```
/* Reader */  
wait(S);  
readcount++;  
if (readcount == 1)  
    wait(wrt);  
signal(S);  
...  
reading is performed  
...  
wait(S);  
readcount--;  
if (readcount == 0)  
    signal(wrt);  
signal(S);
```

The First Readers-Writers Problem

➤ The first readers-writers problem



The Second Readers-Writers Problem

- The second readers-writers problem
 - Give preferential treatment to writers
 - Readers may suffer unbounded waiting

- Shared data
 - semaphore $S1, S2, wrt, rd, wrtpending$;
 - Initially,
 - $S1 = 1, S2 = 1, wrt = 1, rd = 1, wrtpending = 1,$
 - $readcount = 0, writecount = 0$

The Second Readers-Writers Problem

```

/* Writer */
wait(S2);
writecount++;
if (writecount == 1)
    wait(rd);
signal(S2);
wait(wrt)
    ...
    writing is performed
    ...
signal(wrt);
wait(S2);
writecount--;
if (writecount == 0)
    signal(rd);
signal(S2);

```

```

/* Reader */
wait(wrtpending);
wait(rd);
wait(S1);
readcount++;
if (readcount == 1)
    wait(wrt);
signal(S1);
signal(rd);
signal(wrtpending);
    ...
    reading is performed
    ...
wait(S1);
readcount--;
if (readcount == 0)
    signal(wrt);
signal(S1);

```

The Second Readers-Writers Problem

