

# Multicore Operating Systems

Minsoo Ryu

Department of Computer Science and Engineering  
Hanyang University



# Outline

1 Kernels for Small Multiprocessors

Page X

2 Advantages and Limitations of SMP Kernels

Page X

3 Kernels for Large Multiprocessors

Page X

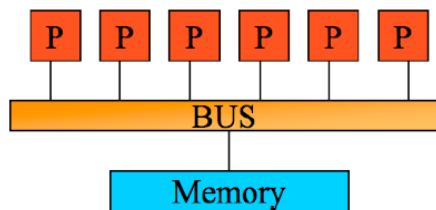
4 Q & A

Page X

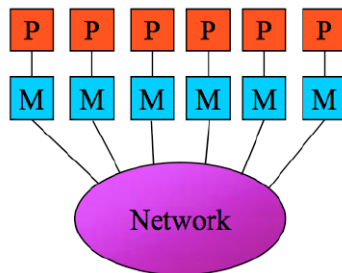
# Considerations for Multiprocessors

## ➤ Diversity in HW organization

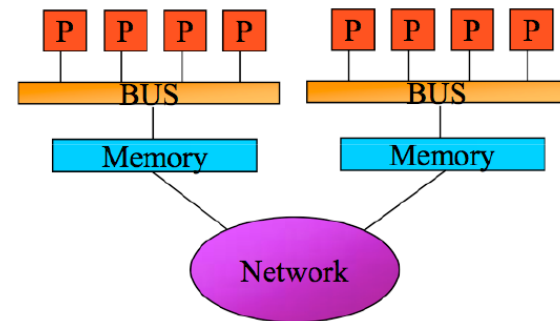
- Shared memory vs. distributed memory
- Bus-based vs. network-based
- Homogeneous vs. heterogeneous



Shared memory &  
bus-based (UMA)



Distributed memory &  
network-based (NUMA)



Clustered & hybrid  
(NUMA)

## ➤ Interactions among processors

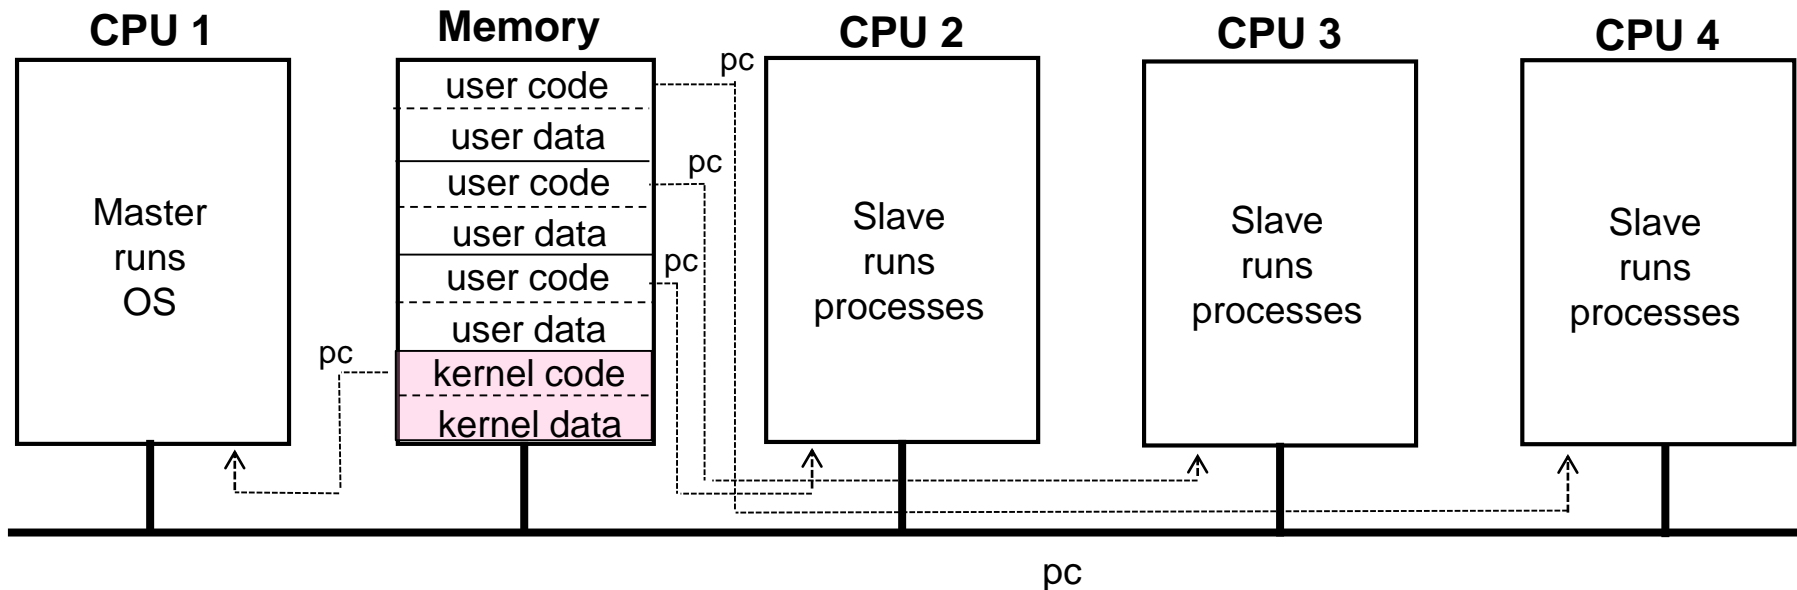
- Memory and cache sharing
- Inter-processor communications

# Classification of Multiprocessor Kernels

- For small multiprocessors
  - Master-slave organization
  - Independent kernels
  - SMP (Symmetric Multiprocessing) kernel
  
- For large multiprocessors (more than tens of cores)
  - Distributed kernels
  - Factored kernels
  - Clustered organization

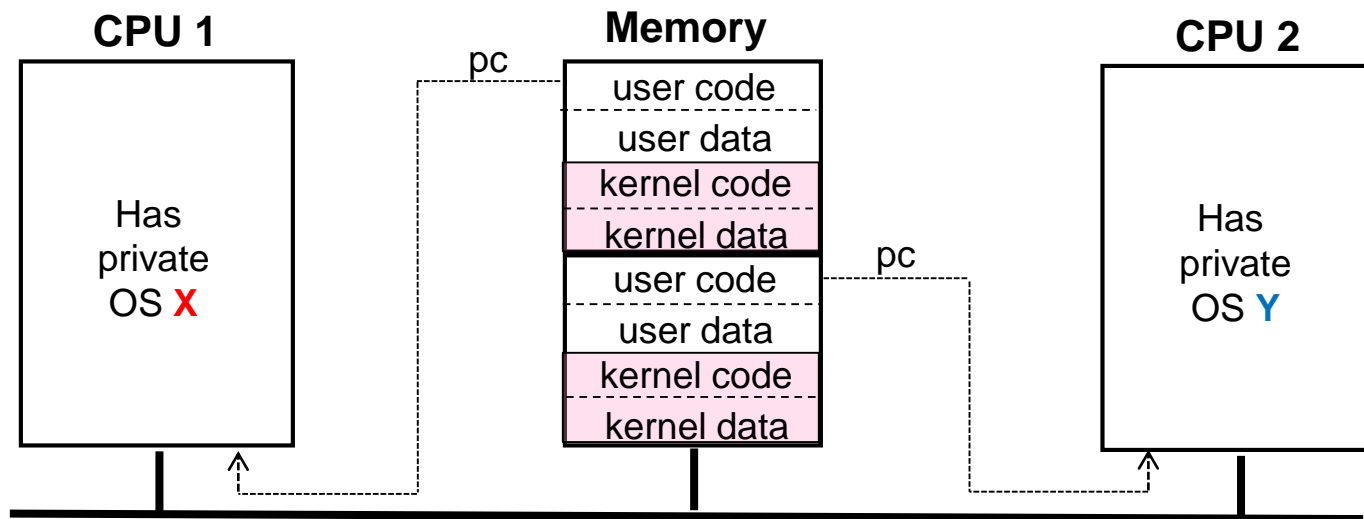
# Master-Slave Organization

- One copy of the OS on a master CPU
  - All system calls are redirected to the master CPU
  - All the other CPUs run user processes
- Problem
  - The master CPU can be a bottleneck



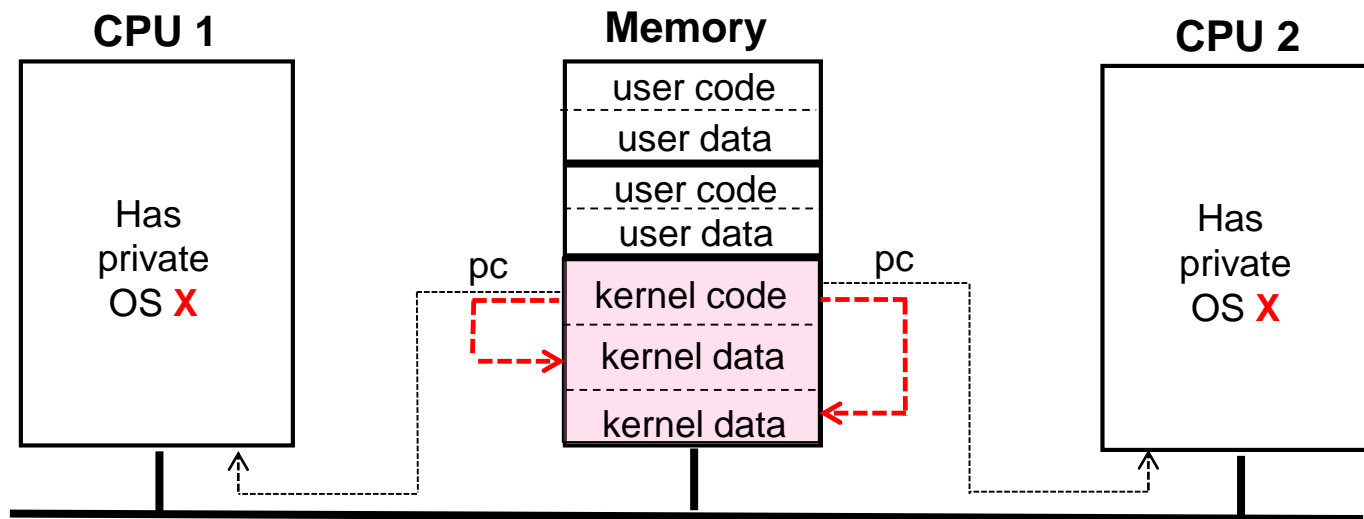
# Independent Kernels

- Each CPU has its own private OS
  - Memory is partitioned and assigned to each OS as private memory
  - The  $n$  CPUs operate as  $n$  independent computers



# Independent Kernel Instances

- Each CPU has a different instance of the same OS
  - All the CPUs share the OS code and make private copies of only the data



# Advantages and Disadvantages

## ➤ Advantages

- Better than  $n$  separate computers since it allows all the machines to share a set of hardware resources
- A single kernel crash does not affect other kernels

## ➤ Disadvantages

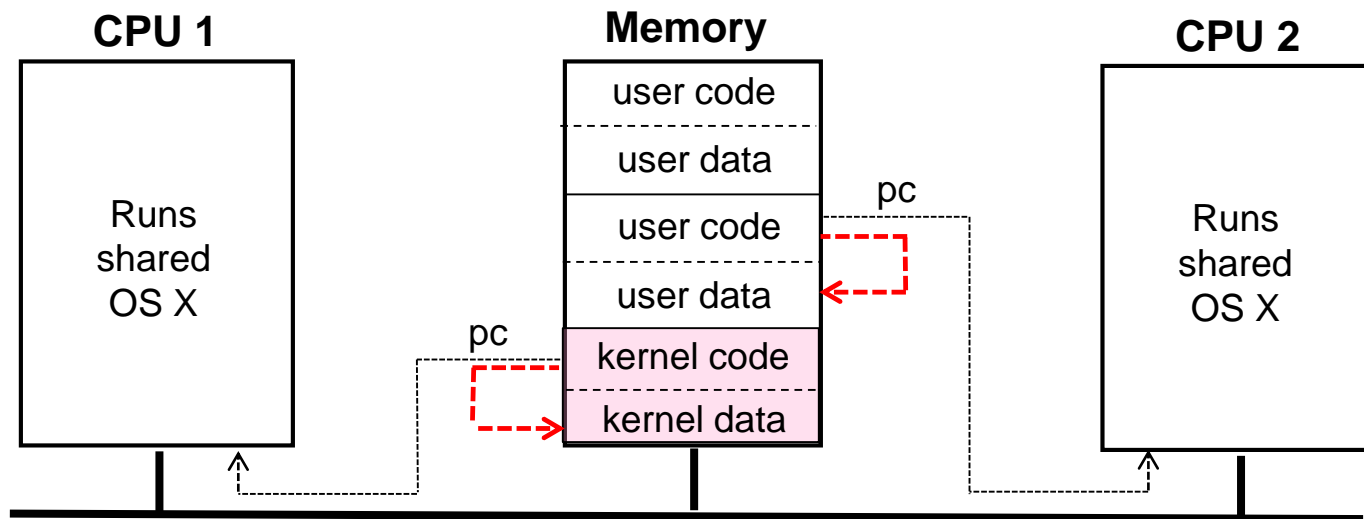
- Applications need to be written in a distributed fashion
  - Cannot exploit the tightly coupled hardware architecture
- There is no sharing of threads and each OS schedules their threads by itself
- Using shared resources may result in inconsistent results
  - e.g.) A certain block is present and dirty in multiple buffer caches at the same time



# SMP (Symmetric Multiprocessing) Kernels

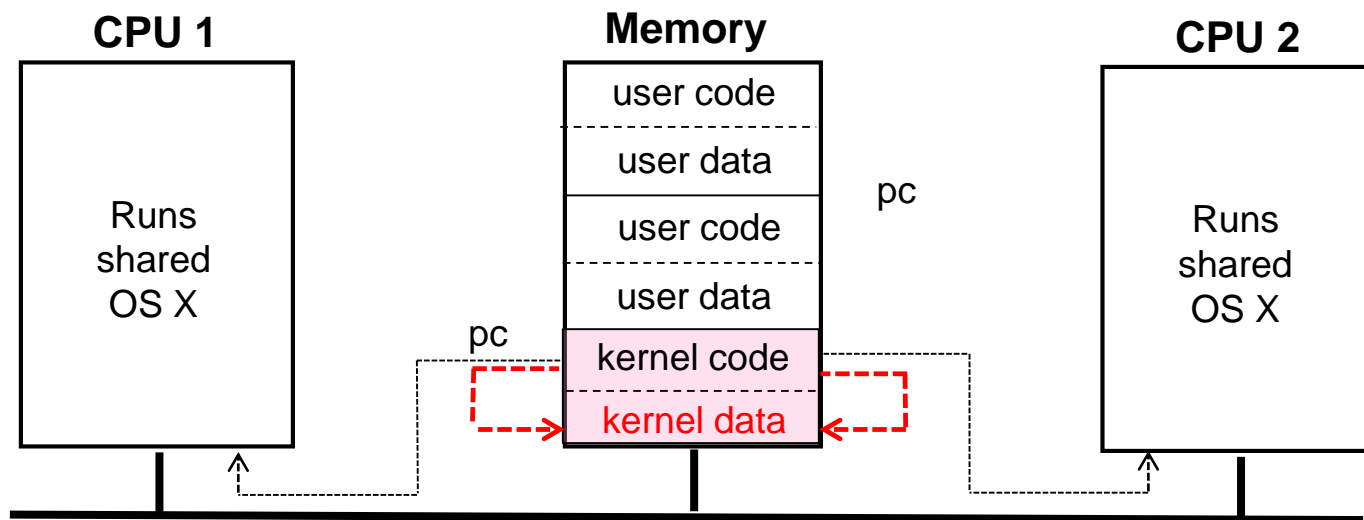
가장 많이 쓰이는 형태의 kernel 형태.

- One copy of the OS in memory and any CPU can run it
  - When a system call is made, the CPU on which the system call was made traps to the kernel and processes the system call



# Data Races

- However, there can be data races
- Two CPUs can execute simultaneously in the kernel
  - They may access the same kernel data simultaneously



# Giant Lock

## ➤ Simple solution

- Associate a mutex (giant lock) with the OS to make the whole system one big critical region
- When a CPU wants to run OS, it must acquire the mutex

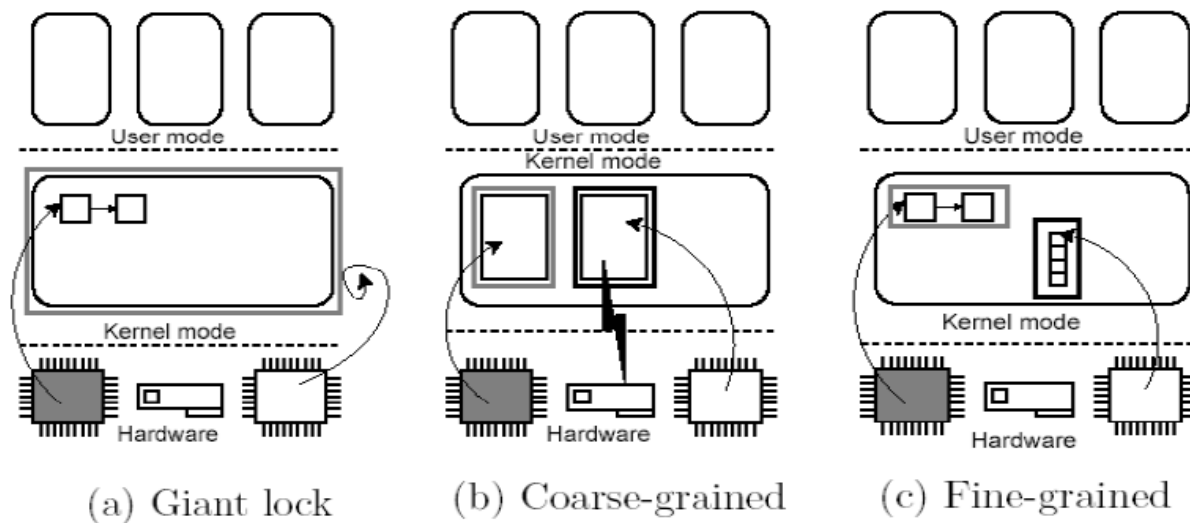
## ➤ This solution is not scalable

- If 10% of all run time is spent inside the OS, then 10 CPUs will pretty much saturate the whole system

# Coarse- and Fine-Grained Locks

## ➤ Better solution

- Split the OS into independent critical regions
  - One CPU can run the scheduler while another CPU is handling a system call
  - Note: great care must be taken to avoid deadlocks



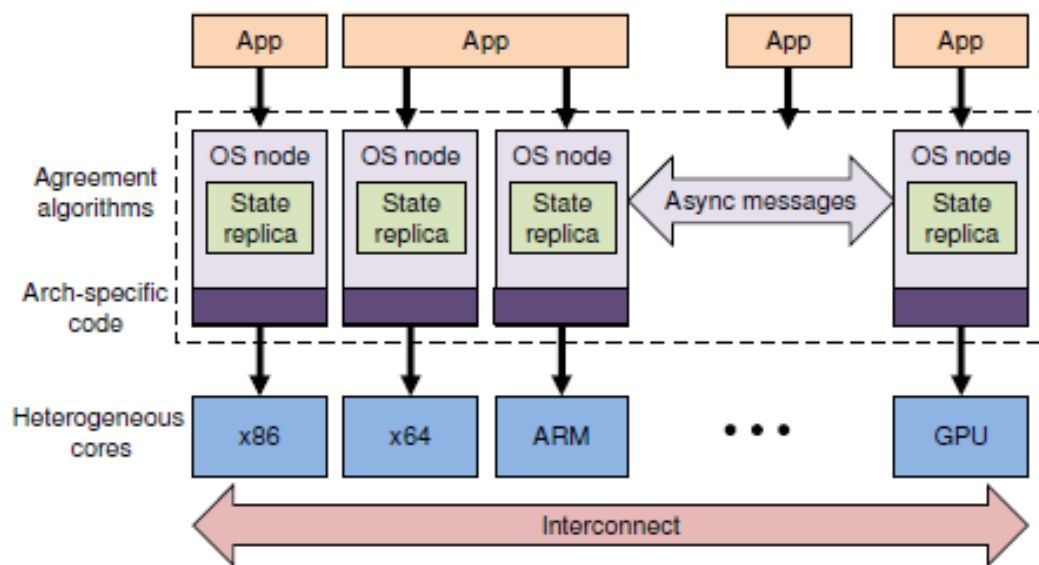
SMP kernels are in widespread use: Unix, Linux, MS Windows, Apple OS X, iOS, ...

# Advantages and Limitations of SMP kernels

- **Advantages: “shared everything structure”**
  - **All OS services and abstractions are shared**
    - e.g.) abstraction of a single flat, global address space
  - **All HW resources are shared**
  
- **Limitations**
  - **SMP kernels require restricted (SMP) HW architectures**
    - Identical instruction sets across processors
    - Shared memory structure
  - **SMP kernels do not scale well**
    - Lock contention cost
    - Reliance on shared memory

# Distributed Kernels: The Multikernel Model

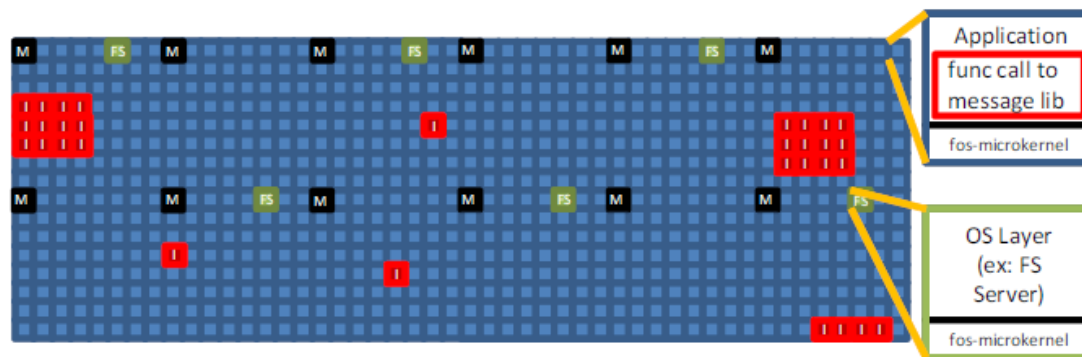
- Use ideas from distributed systems
  - Treats the machine as a network of independent cores
  - Replicates OS states across cores that communicate using messages and share no memory



Microsoft Research and ETH Zurich,  
ACM SOSP, 2009

# Factored Kernels: FOS

- A new operating system targeting manycore systems
  - Each operating system service is factored into a set of communicating servers
    - Space sharing rather than time sharing
  - Relies on message passing rather shared memory



Space multiplexing

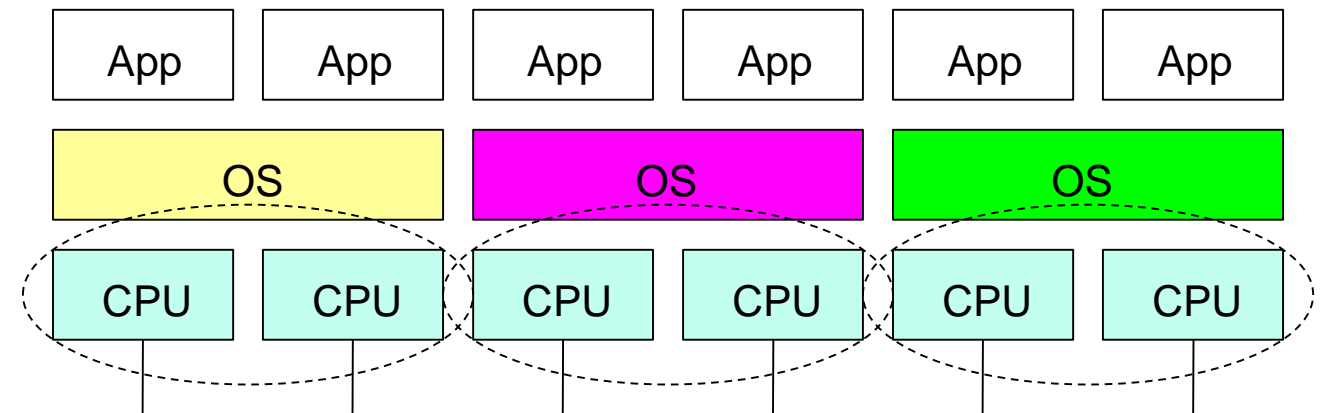
MIT, ACM OSR, 2009

- I - Idle Processor Core
- - Application
- FS FS ... FS - Fleet of File System Servers
- M M ... M - Fleet of Physical Memory Allocation Servers

# Clustered Organization

## ➤ Key idea

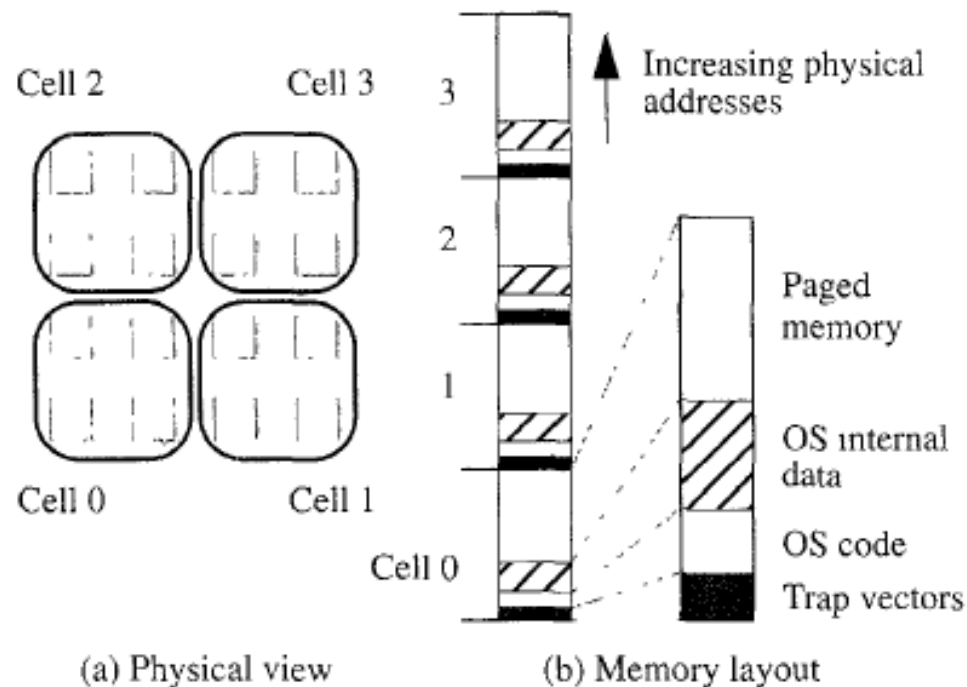
- Processors are grouped into clusters
- Each cluster runs an independent OS kernel
- Clusters are the unit of scalability





# Hive (Stanford, ACM SOSP'95)

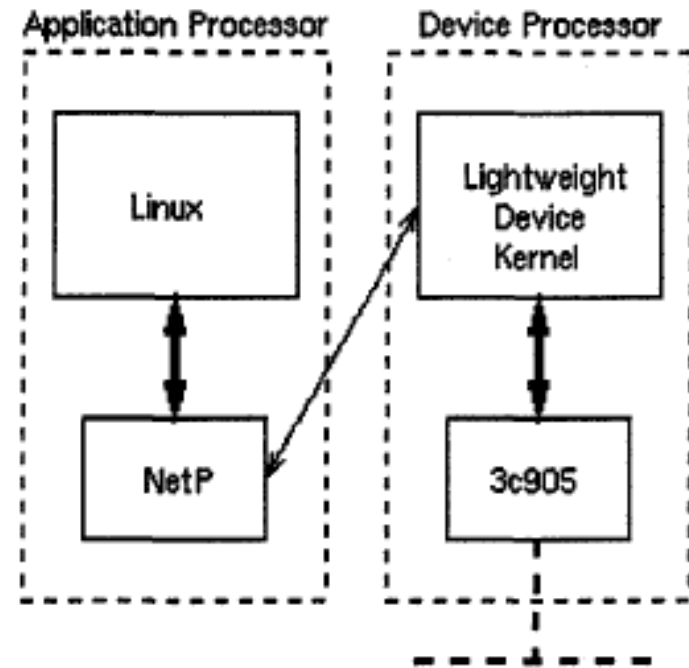
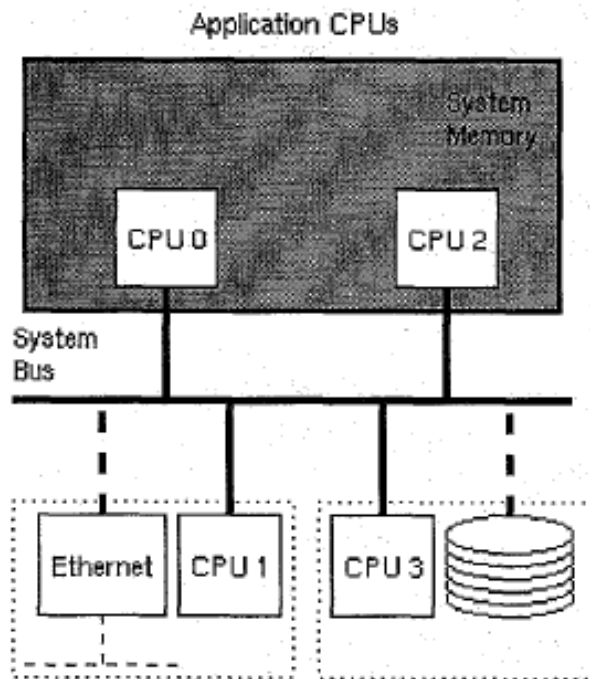
- The original purpose of Hive was to achieve reliability and fault containment
- Each cell runs an SMP kernel



# AsyMOS: Asymmetric Independent Kernels

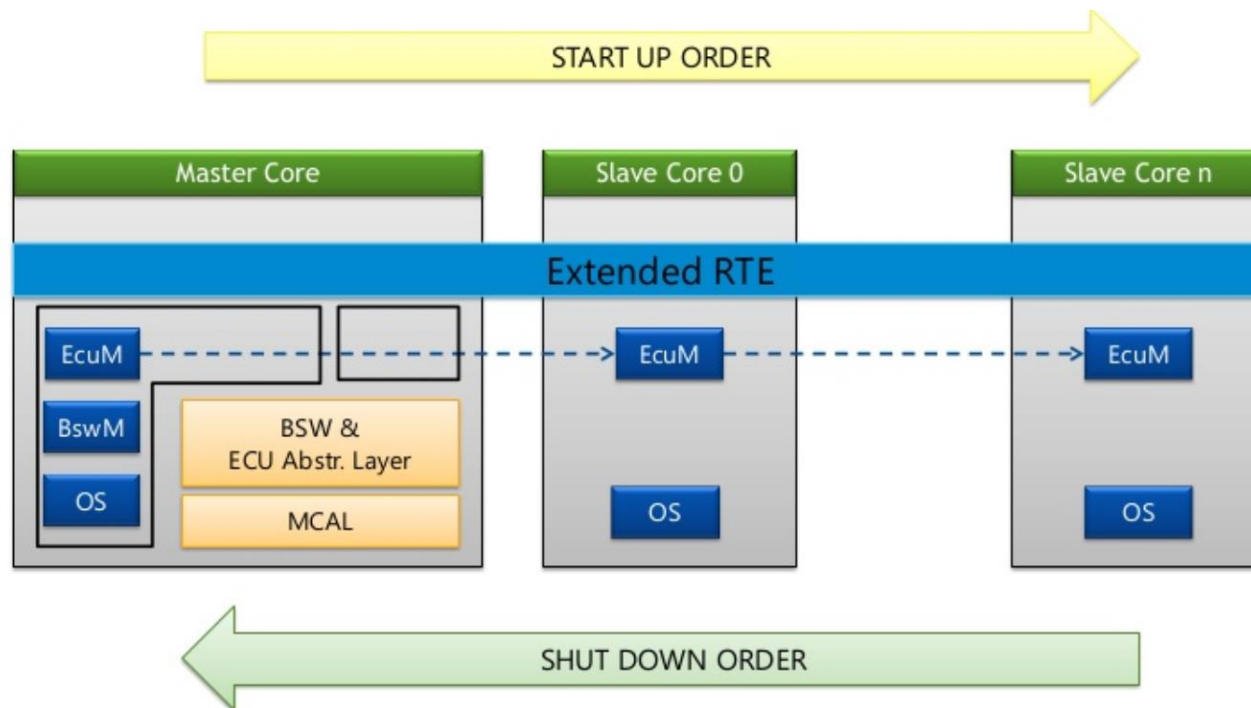
## ➤ Two types of kernel

- A native OS kernel running on application processors
- LDK (lightweight device kernel) running on device processors



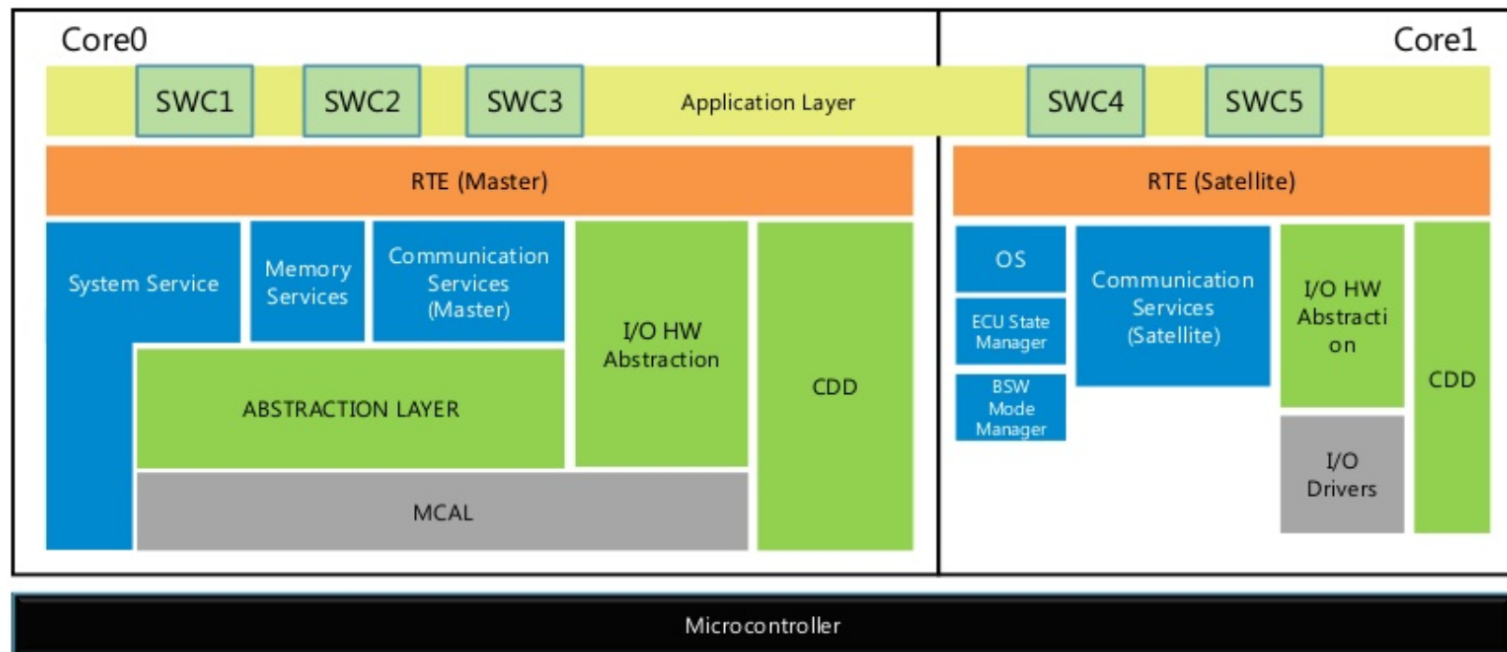
# AUTOSAR Multicore: Master-Slave Concept

- Static assignment of OS-Applications to cores
- Master core contains complete BSW
- Master core controls start up and shut down



# Master-Satellite Concept

- The partitioning is implementation specific and the communication between master and satellite is not standardized
  - BSW can be partitioned (eg., a FlexRay cluster is on one core and a CAN cluster on a difference core)



# Spinlocks and IOC

## ➤ Spinlocks

- protects a variable across core so that no two cores may access it at the same time
- Is a busy waiting mechanism

## ➤ IOC (Inter OS-Application Communication)

- Provides communication services which can be accessed by clients which need to communicate across cores and memory protection boundaries

