

Multicore Operating Systems

Minsoo Ryu

Department of Computer Science and Engineering
Hanyang University



Outline

1 Kernels for Small Multiprocessors

Page X

2 Advantages and Limitations of SMP Kernels

Page X

3 Kernels for Large Multiprocessors

Page X


4 Q & A

Page X


Considerations for Multiprocessors

➤ Diversity in HW organization


- Shared memory vs. distributed memory
- Bus-based vs. network-based
- Homogeneous vs. heterogeneous



Shared memory &
bus-based (UMA)



Distributed memory &
network-based (NUMA)



Clustered & hybrid
(NUMA)

➤ Interactions among processors

- Memory and cache sharing
- Inter-processor communications

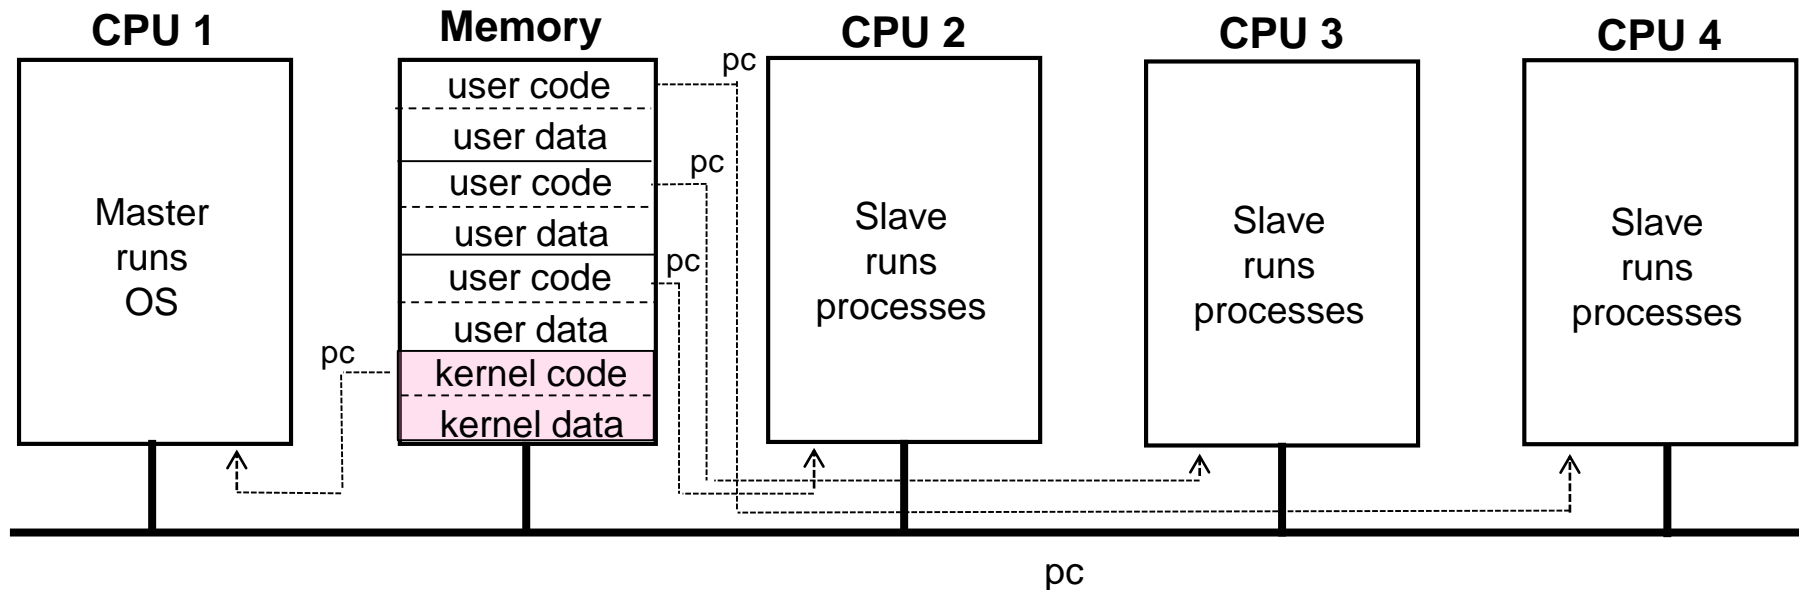
Classification of Multiprocessor Kernels

- For small multiprocessors
 - Master-slave organization
 - Independent kernels
 - SMP (Symmetric Multiprocessing) kernel

- For large multiprocessors (more than tens of cores)
 - Distributed kernels
 - Factored kernels
 - Clustered organization

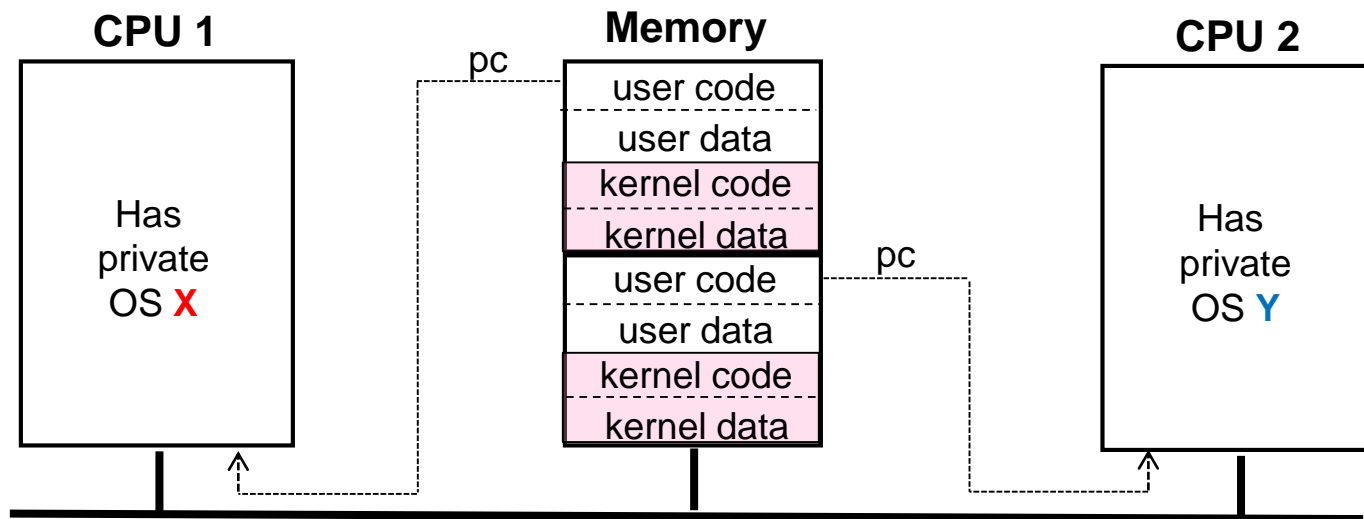
Master-Slave Organization

- One copy of the OS on a master CPU
 - All system calls are redirected to the master CPU
 - All the other CPUs run user processes
- Problem
 - The master CPU can be a bottleneck



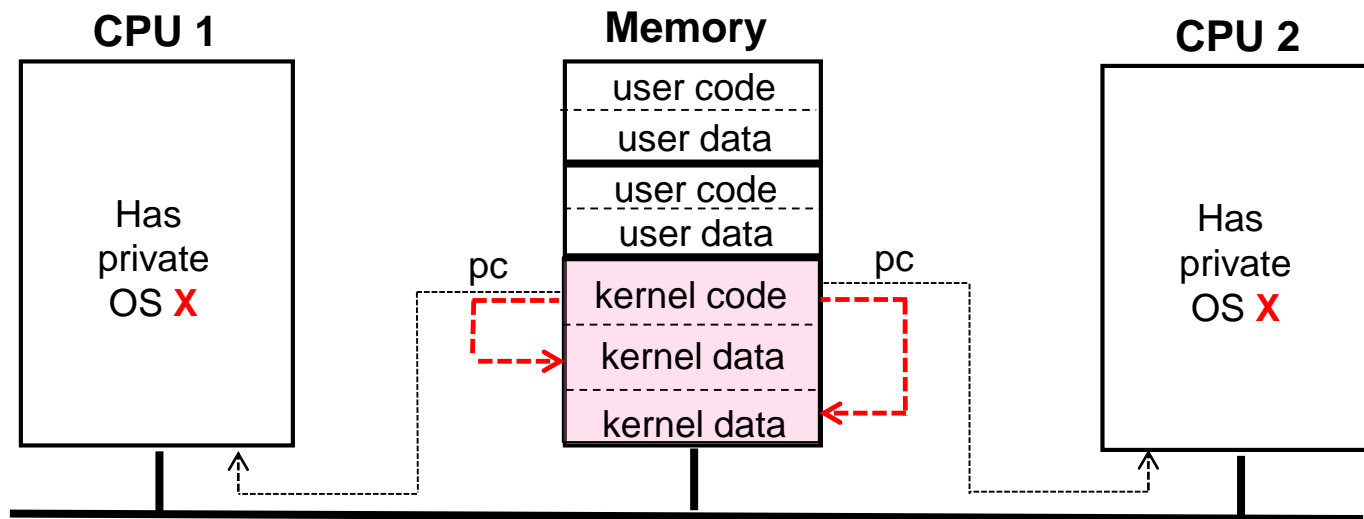
Independent Kernels

- Each CPU has its own private OS
 - Memory is partitioned and assigned to each OS as private memory
 - The n CPUs operate as n independent computers



Independent Kernel Instances

- Each CPU has a different instance of the same OS
 - All the CPUs share the OS code and make private copies of only the data



Advantages and Disadvantages

➤ Advantages

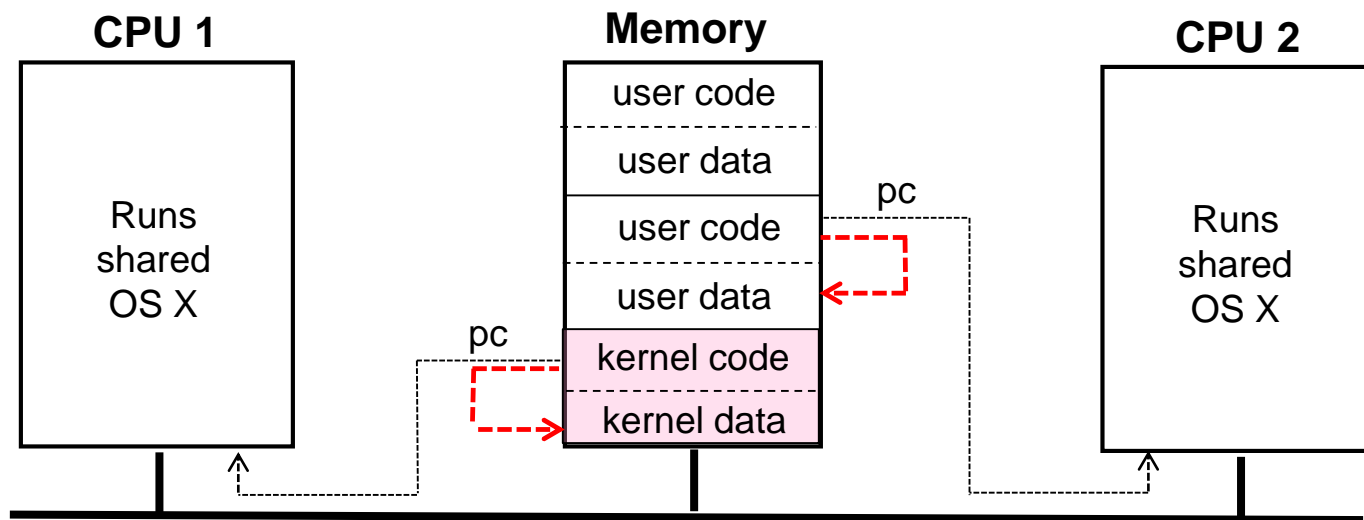
- Better than n separate computers since it allows all the machines to share a set of hardware resources
- A single kernel crash does not affect other kernels

➤ Disadvantages

- Applications need to be written in a distributed fashion
 - Cannot exploit the tightly coupled hardware architecture
- There is no sharing of threads and each OS schedules their threads by itself
- Using shared resources may result in inconsistent results
 - e.g.) A certain block is present and dirty in multiple buffer caches at the same time

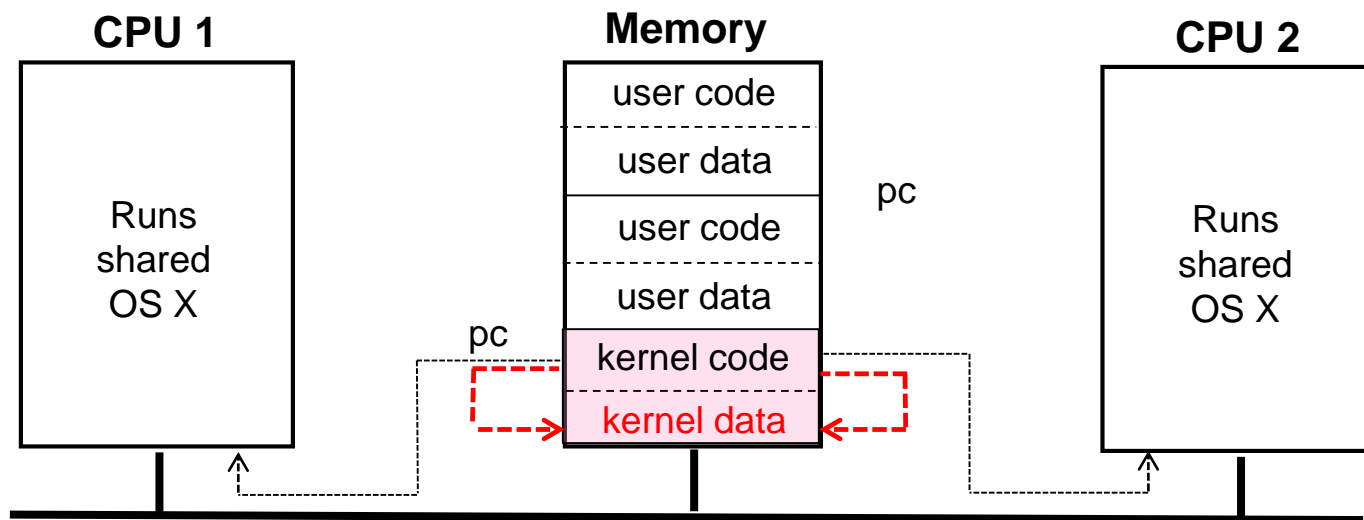
SMP (Symmetric Multiprocessing) Kernels

- One copy of the OS in memory and any CPU can run it
 - When a system call is made, the CPU on which the system call was made traps to the kernel and processes the system call



Data Races

- However, there can be data races
- Two CPUs can execute simultaneously in the kernel
 - They may access the same kernel data simultaneously



Giant Lock

➤ Simple solution

- Associate a mutex (giant lock) with the OS to make the whole system one big critical region
- When a CPU wants to run OS, it must acquire the mutex

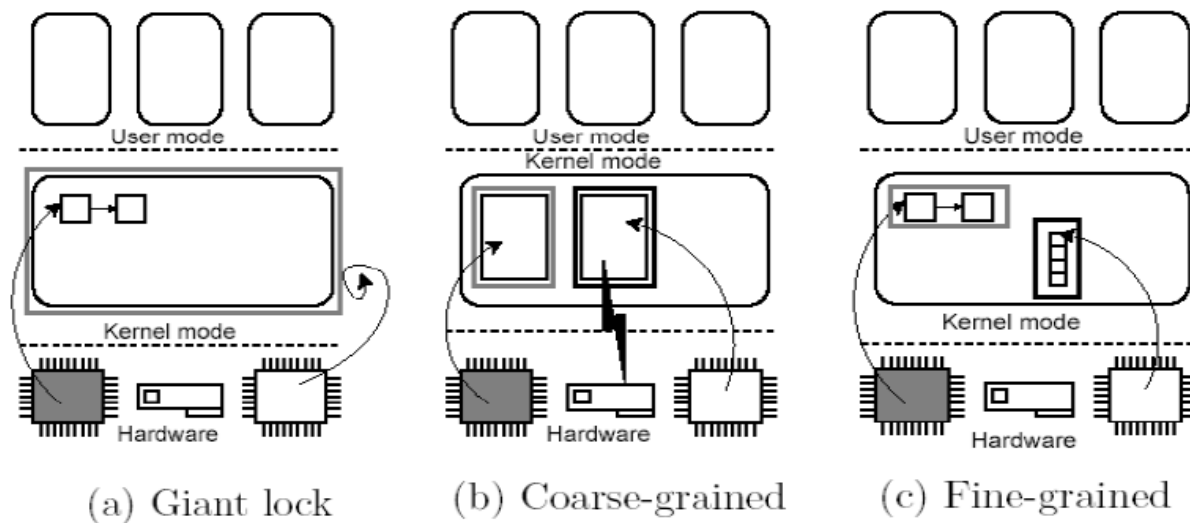
➤ This solution is not scalable

- If 10% of all run time is spent inside the OS, then 10 CPUs will pretty much saturate the whole system

Coarse- and Fine-Grained Locks

➤ Better solution

- Split the OS into independent critical regions
 - One CPU can run the scheduler while another CPU is handling a system call
 - Note: great care must be taken to avoid deadlocks



SMP kernels are in widespread use: Unix, Linux, MS Windows, Apple OS X, iOS, ...

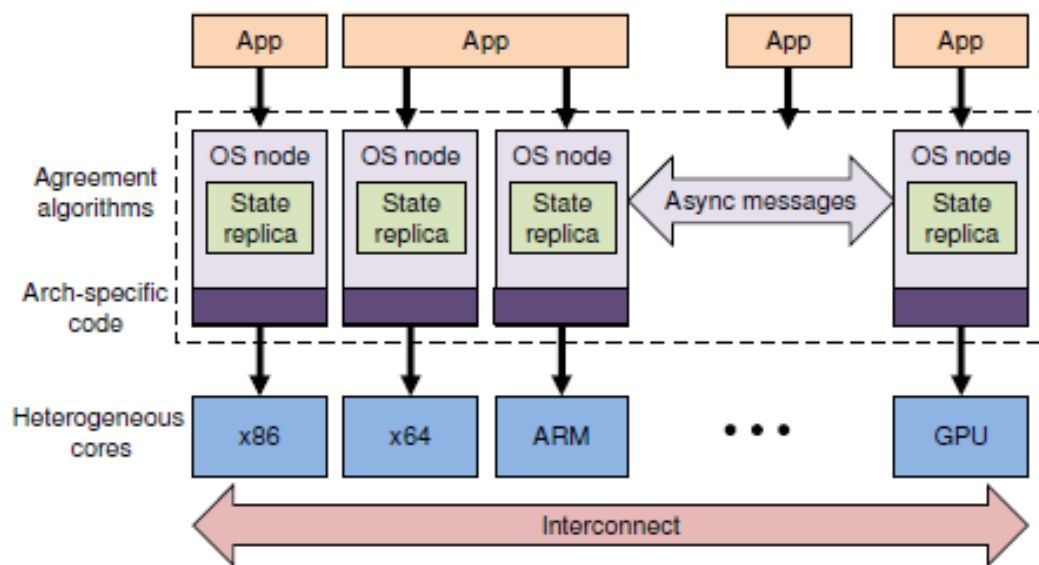
Advantages and Limitations of SMP kernels

- **Advantages: “shared everything structure”**
 - **All OS services and abstractions are shared**
 - e.g.) abstraction of a single flat, global address space
 - **All HW resources are shared**

- **Limitations**
 - **SMP kernels require restricted (SMP) HW architectures**
 - Identical instruction sets across processors
 - Shared memory structure
 - **SMP kernels do not scale well**
 - Lock contention cost
 - Reliance on shared memory

Distributed Kernels: The Multikernel Model

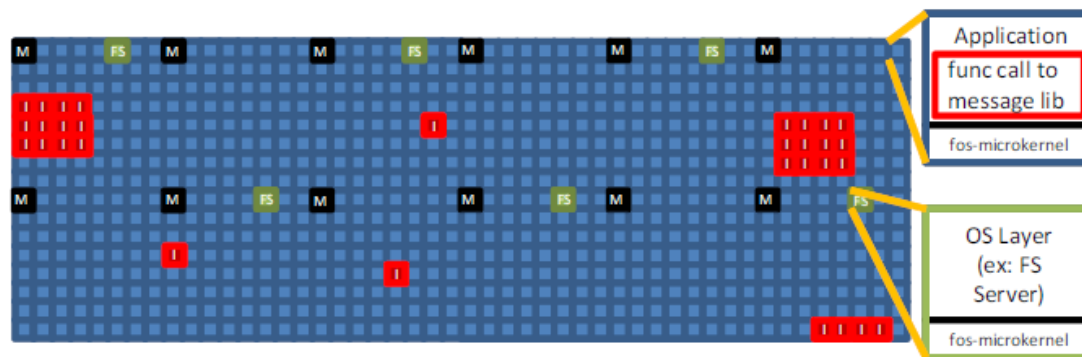
- Use ideas from distributed systems
 - Treats the machine as a network of independent cores
 - Replicates OS states across cores that communicate using messages and share no memory



Microsoft Research and ETH Zurich,
ACM SOSP, 2009

Factored Kernels: FOS

- A new operating system targeting manycore systems
 - Each operating system service is factored into a set of communicating servers
 - Space sharing rather than time sharing
 - Relies on message passing rather shared memory



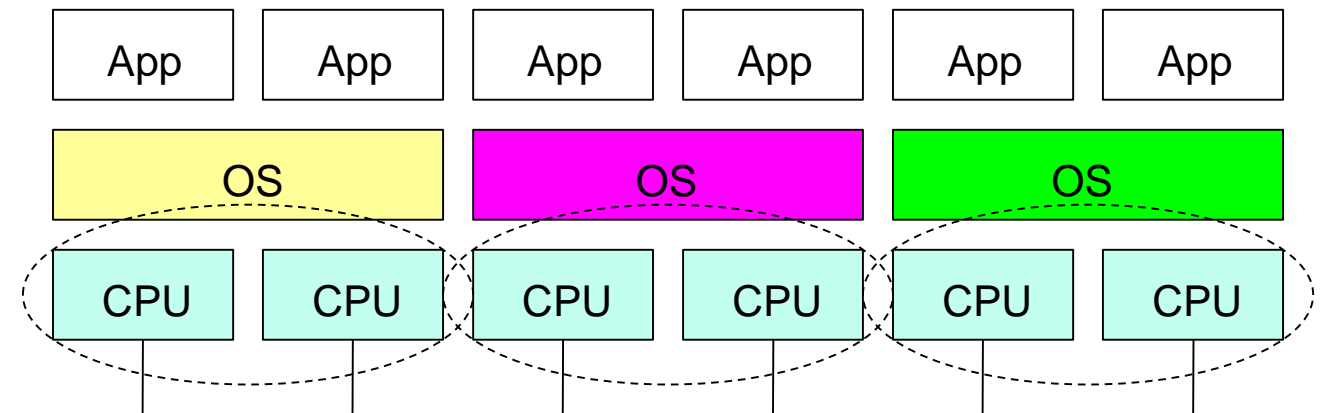
MIT, ACM OSR, 2009

- - Idle Processor Core
- - Application
- FS FS ... FS - Fleet of File System Servers
- M M ... M - Fleet of Physical Memory Allocation Servers

Clustered Organization

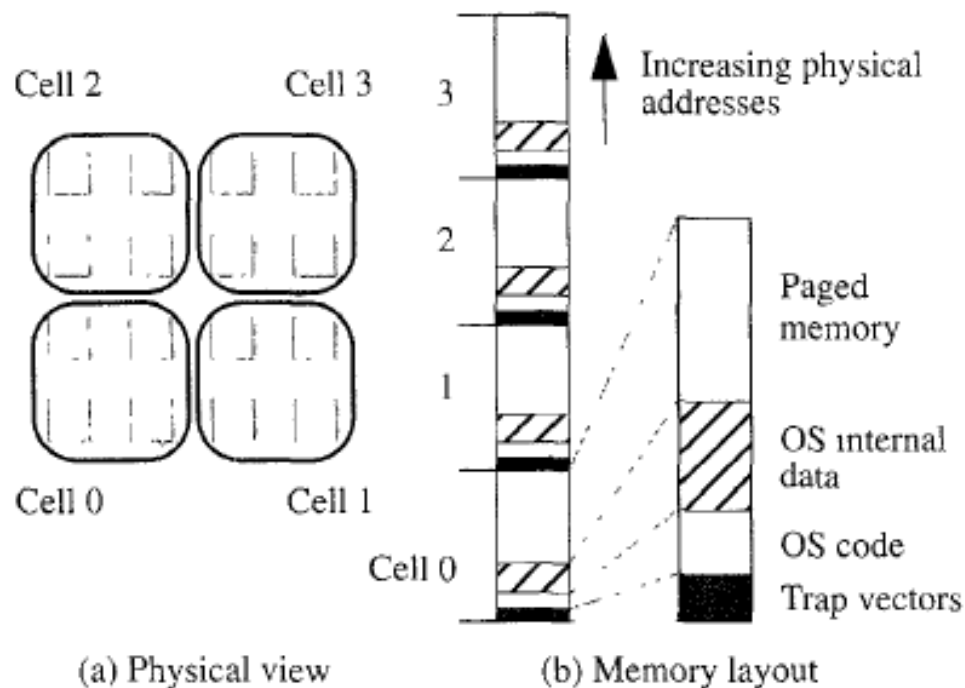
➤ Key idea

- Processors are grouped into clusters
- Each cluster runs an independent OS kernel
- Clusters are the unit of scalability



Hive (Stanford, ACM SOSP'95)

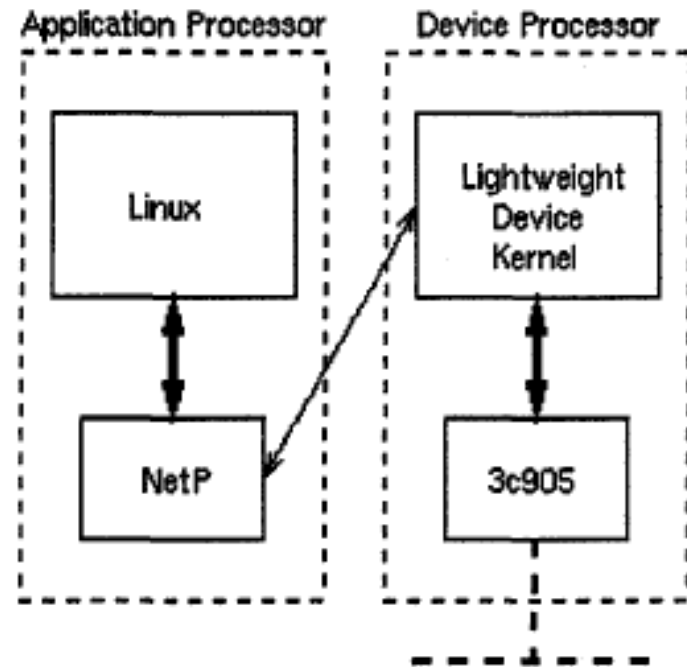

- The original purpose of Hive was to achieve reliability and fault containment
- Each cell runs an SMP kernel



AsyMOS: Asymmetric Independent Kernels

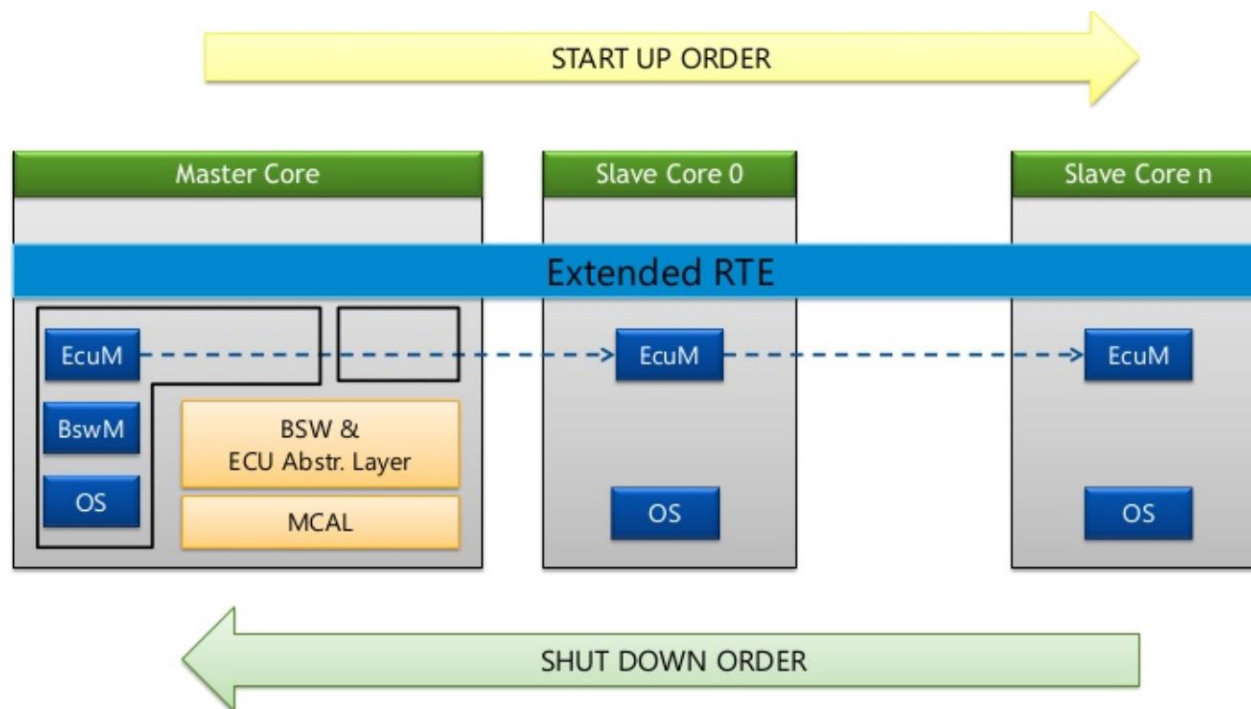
➤ Two types of kernel

- A native OS kernel running on application processors
- LDK (lightweight device kernel) running on device processors



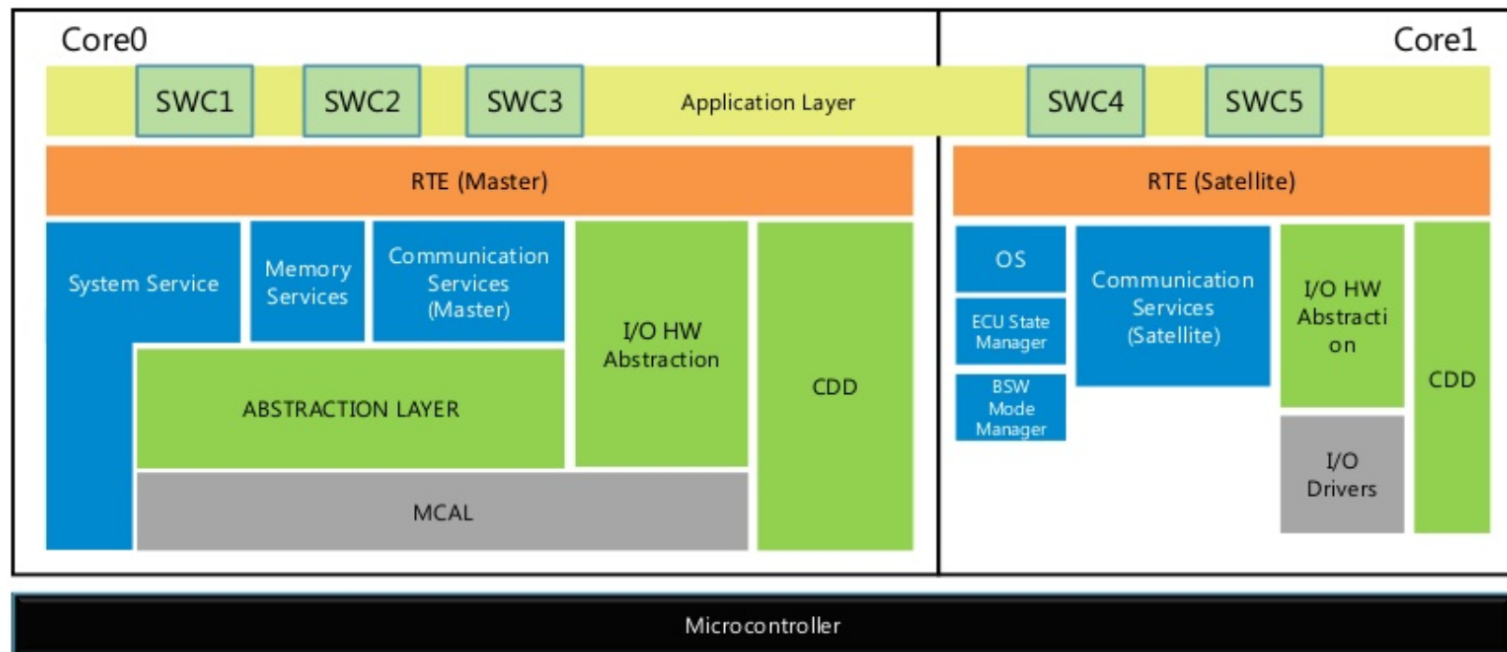
AUTOSAR Multicore: Master-Slave Concept

- Static assignment of OS-Applications to cores
- Master core contains complete BSW
- Master core controls start up and shut down



Master-Satellite Concept

- The partitioning is implementation specific and the communication between master and satellite is not standardized
 - BSW can be partitioned (eg., a FlexRay cluster is on one core and a CAN cluster on a difference core)



Spinlocks and IOC

➤ Spinlocks

- protects a variable across core so that no two cores may access it at the same time
- Is a busy waiting mechanism

➤ IOC (Inter OS-Application Communication)

- Provides communication services which can be accessed by clients which need to communicate across cores and memory protection boundaries

