

Multicore Performance #1

- CPU Scheduling -

Minsoo Ryu

Department of Computer Science and Engineering
Hanyang University



Outline

1

CPU Scheduling*Page***X**

2

Priority-Based Scheduling*Page***X**

3

Proportional Share Scheduling*Page***X**

4

Multiprocessor Scheduling*Page***X**

CPU Scheduling

- CPU scheduling is to decide when and which process to run among ready processes

- Scheduling criteria
 - Performance objectives
 - Maximize CPU utilization or throughput
 - Minimize completion time, waiting time, or response time
 - Real-time constraints
 - Satisfy deadlines
 - Fairness
 - Provide CPU cycles proportional to weights

Classification of Scheduling Policies

- **Two paradigms**
 - Priority-based scheduling
 - Proportional share scheduling

- **Number of processors**
 - Uniprocessor scheduling
 - Multiprocessor scheduling

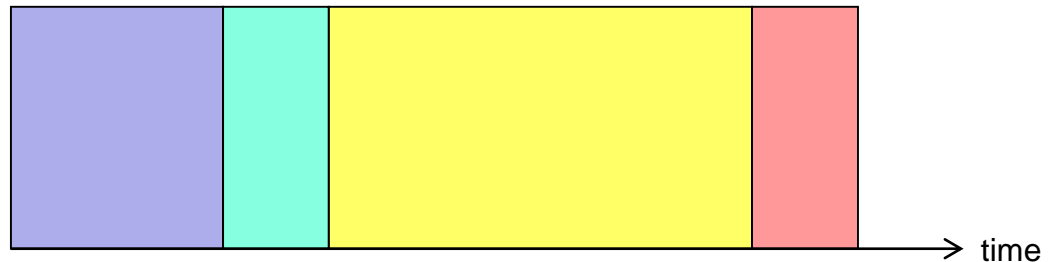
Priority-Based Scheduling



Priority-Based Scheduling

- A priority number (integer) is associated with each process

- Task A (priority 1)
- Task B (priority 2)
- Task C (priority 3)
- Task D (priority 4)



- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive

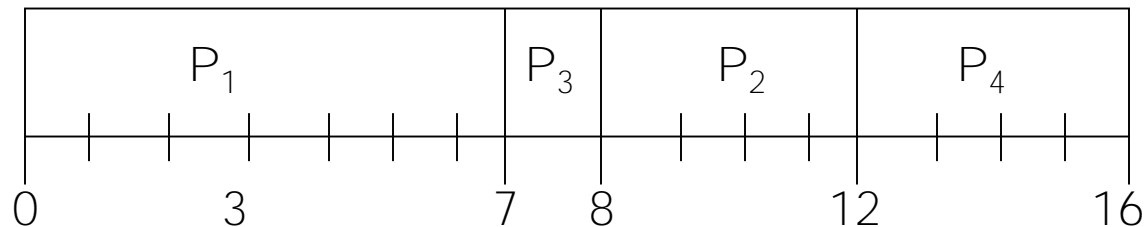
Priority-Based Scheduling Policies

- **Non-real-time policies**
 - FCFS (First-Come, First-Served)
 - SJF (Shortest-Job-First)
 - SRTF (Shortest-Remaining-Time-First)

- **Real-time policies**
 - RM (Rate Monotonic)
 - EDF (Earliest Deadline First)

SJF (Shortest-Job-First)

<u>Process</u>	<u>Arrival Time</u>	<u>Execution Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4



➤ Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

SRTF (Shortest-Remaining-Time-First)


<u>Process</u>	<u>Arrival Time</u>	<u>Execution Time</u>
----------------	---------------------	-----------------------

P_1	0.0	7
-------	-----	---

P_2	2.0	4
-------	-----	---

P_3	4.0	1
-------	-----	---

P_4	5.0	4
-------	-----	---



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$
 - SRTF is optimal in that it achieves minimum average waiting time for a given set of processes

RM (Rate Monotonic)

빈도가 높은 일에 우선순위를 주는 방식.

➤ Assumptions

- Processes have periods, worst-case execution times (WCETs), and deadlines

➤ Scheduling policy

- Give higher priorities to tasks with shorter periods
- Preemptive static priority scheduling

➤ Optimality

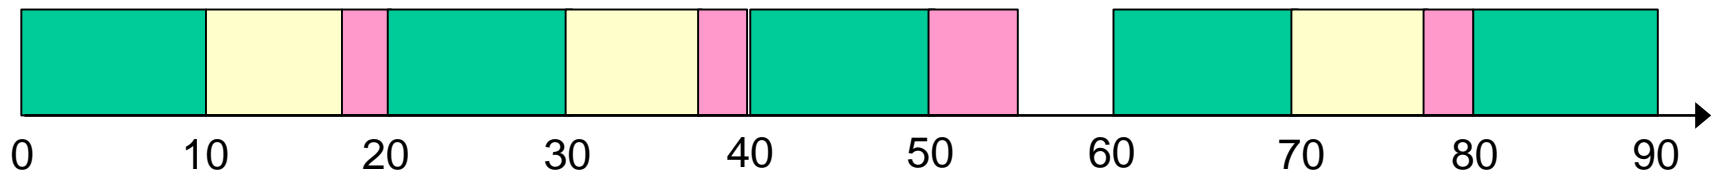
- If a feasible static priority assignment exists for some process set, the RM priority assignment is feasible for that process set

RM (Rate Monotonic)

➤ Consider the following tasks

- Process X ■: period = 20, WCET = 10, deadline = 20
- Process Y ■: period = 30, WCET = 8, deadline = 30
- Process Z ■: period = 40, WCET = 4, deadline = 40

period : 주기



➤ Schedulability test

m : CPU count
U : CPU utilization

cpu utilization이 69%를 넘지 않으면
rate monotonic을 할 수 있다.

- CPU utilization: $U = \sum_{i=1}^m e_i / p_i$
- A set of m processes is schedulable if $U \leq m(2^{1/m} - 1)$
 - For large m, $m(2^{1/m} - 1) \approx \ln 2 \approx 0.69$

EDF (Earliest Deadline First)

➤ Scheduling policy

- Give higher priorities to tasks with earlier absolute deadlines
- Preemptive **dynamic priority scheduling**

➤ Optimality

- If a feasible dynamic priority assignment exists for some process set, the EDF priority assignment is feasible for that process set

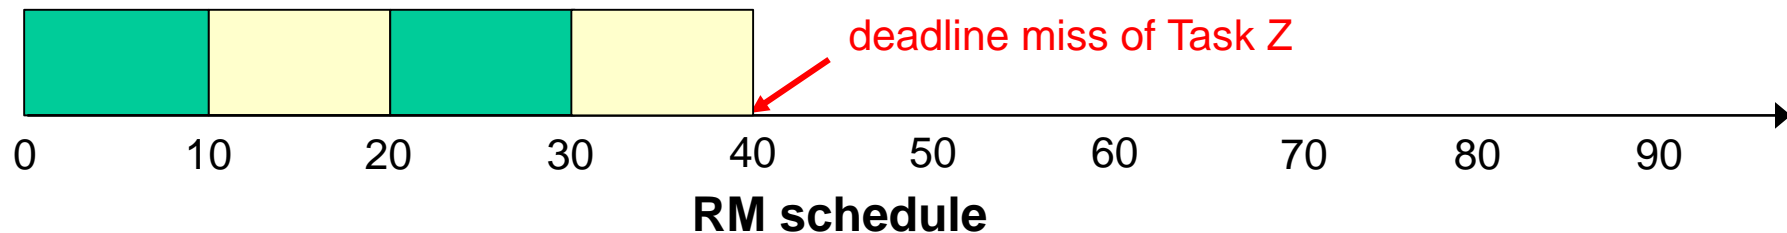
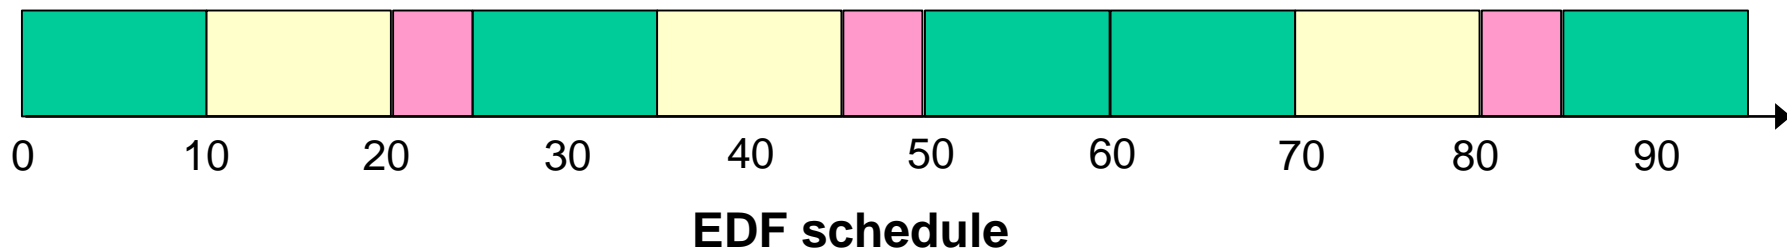
➤ Schedulability test

- A set of m processes is schedulable if and only if $U \leq 1$

EDF (Earliest Deadline First)




➤ Consider the following tasks

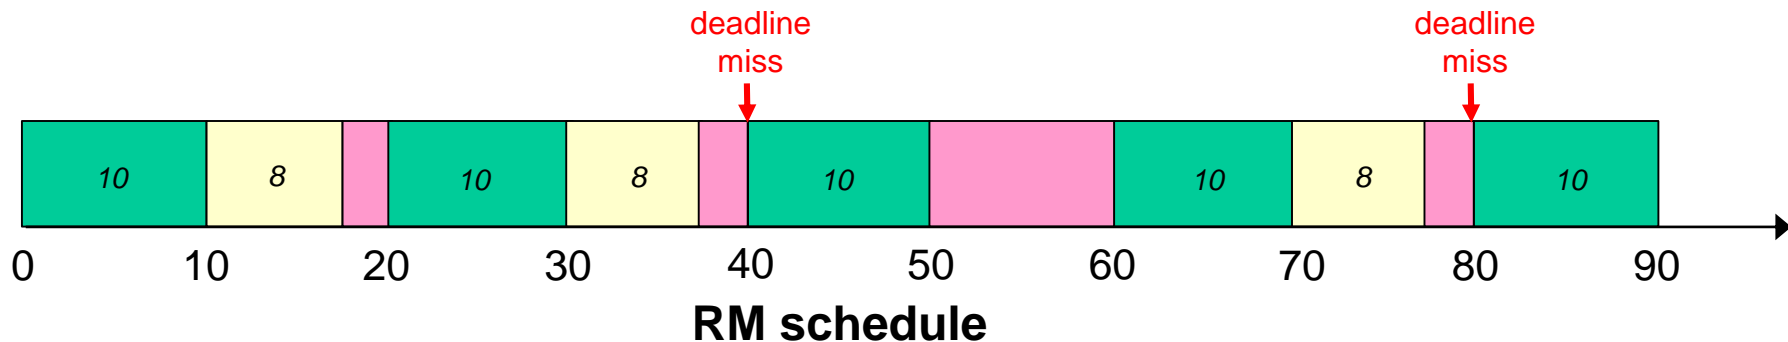
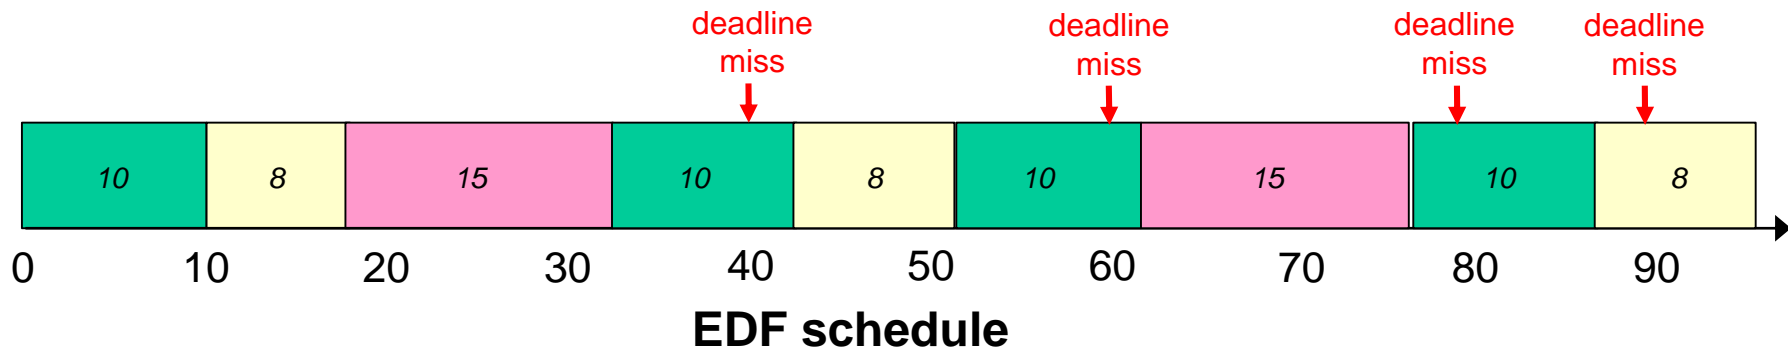
- Process X ■: period = 20, WCET = 10, deadline = 20
- Process Y ■: period = 30, WCET = 10, deadline = 30
- Process Z ■: period = 40, WCET = 5, deadline = 40



Non-schedulable Behavior

➤ Consider the following tasks

- Process X : period = 20, WCET = 10, deadline = 20
- Process Y : period = 30, WCET = 8, deadline = 30
- Process Z : period = 40, WCET = 15, deadline = 40



Proportional Share Scheduling



Proportional Share Scheduling

➤ Basic concept

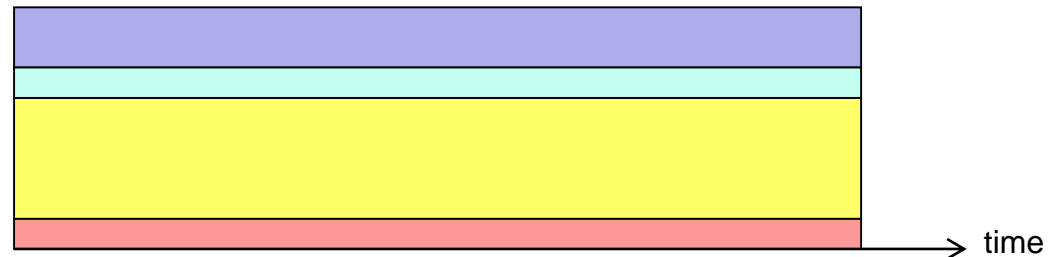
- A weight value is associated with each process
- The CPU is allocated to the process in proportion to its weight

■ Task A (weight 25.0%)

■ Task B (weight 12.5%)

■ Task C (weight 50.0%)

■ Task D (weight 12.5%)



➤ Two contexts

- Fair queueing (in the context of communication networks)
 - Packet scheduling
- Proportional share (in the context of operating systems)
 - Process scheduling

Scheduling Algorithms

➤ Network scheduling

- PGPS (= WFQ), Demers *et al.*, 1989
- Virtual Clock, Lixia Zhang, 1990
- SCFQ, Golestani, 1994
- SFQ, Goyal *et al.*, 1996
- WF²Q, Bennett *et al.*, 1996

➤ CPU scheduling

- Lottery and Stride, Waldspurger, 1995
- Hierarchical SFQ, Goyal *et al.*, 1996
- BVT, Duda *et al.*, 1999
- VTRR, Nieh *et al.*, 2001

GPS (Generalized Processor Sharing)

- A GPS server is defined by (Kleinrock, 1976)

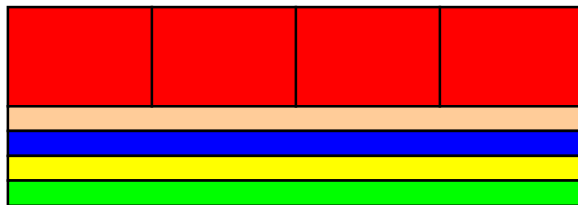
$$\frac{W_i(t_1, t_2)}{W_j(t_1, t_2)} \geq \frac{r_i}{r_j}, j = 1, 2, \dots, N$$

- $W_i(t_1, t_2)$: the amount of session i traffic served in an $(t_1, t_2]$ interval
 - r_i : weight of session i
- Packets of all sessions are served simultaneously
- Idealized fluid-flow system or bit-by-bit weighted round-robin




PGPS (Packet-by-Packet GPS)

- PGPS is an approximation to GPS
 - Fair queueing on a packet-by-packet basis
 - Also known as WFQ (Weighted Fair Queueing)
- Algorithm
 - F_p : the time at which packet p will depart (finish service) under GPS
 - Serve packets in increasing order of F_p



GPS



PGPS

Fairness Bounds of PGPS

➤ Bound on lag

- $\hat{F}_p - F_p \leq \frac{L_{\max}}{r}$
 - \hat{F}_p : the time at which packet p departs under PGPS
 - F_p : the time at which packet p will depart under GPS
 - L_{\max} : the maximum packet length
 - If $r = 1Gbps$ and $L_{\max} = 1Kb$, then the lag is 1 ms

➤ Bound on difference of services received

- $W_i(0, \tau) - \hat{W}_i(0, \tau) \leq L_{\max}$

Worst-case Fair WFQ

➤ WF²Q only considers

- Packets that have started receiving service under GPS

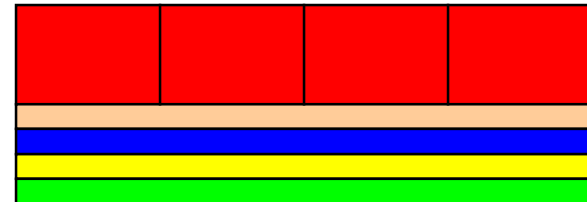
WFQ satisfies only

$$W_p(t_1, t_2) - \hat{W}_p(t_1, t_2) \leq L_{\max}$$

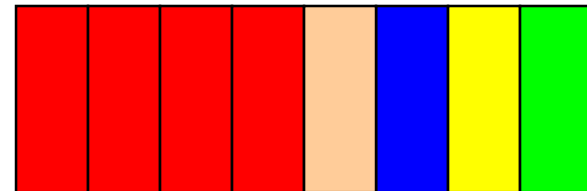
not

$$|W_p(t_1, t_2) - \hat{W}_p(t_1, t_2)| \leq L_{\max}$$

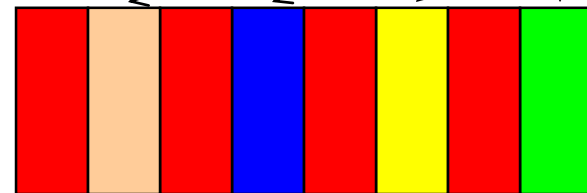
GPS
(fluid-flow)



WFQ



WF²Q



Lottery and Stride Scheduling

- Tickets: relative resource rights
 - Task τ_i has m_i tickets
 - A total of M tickets
- Lottery scheduling: probabilistic algorithm
 - Use random number generator to select a winning ticket
 - Task τ_i is probabilistically guaranteed a rate of $p = m_i / M$
- Stride scheduling: deterministic algorithm
 - “Stride” is inversely proportional to tickets
 - Task with minimum “pass” value is selected and its pass is advanced by its stride

Stride Scheduling

- Task τ_1 : tickets = 3, stride = 2
- Task τ_2 : tickets = 2, stride = 3
- Task τ_3 : tickets = 1, stride = 6

Initial pass values are set to stride values

pass \ time	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
τ_1 's pass value	2	4	4	6	8	8	8	10	10	12	14	14	14	16	16
τ_2 's pass value	3	3	6	6	6	9	9	9	12	12	12	15	15	15	18
τ_3 's pass value	6	6	6	6	6	6	12	12	12	12	12	12	18	18	18

Smallest value is chosen
Pass is advanced by stride 3

Ties are broken arbitrarily

Multiprocessor Scheduling



Uniprocessor vs. Multiprocessor Scheduling

- **Uniprocessor scheduling**
 - It is to decide when and which job will run


- **Multiprocessor scheduling**
 - It is to decide not only when but also where a job will run
 - Almost the same goals as those of uniprocessor scheduling
 - But it raises new issues
 - How to assign applications to multiple processors?
 - How to balance workload among processors?
 - How to define and exploit affinity?
 - How to manage processor heterogeneity?

Multiprocessor Scheduling Policies


- **The same policies as uniprocessor policies**
 - Priority-based scheduling: FCFS, SJF, SRTF, RM, EDF
 - Proportional share scheduling: PGPS, SFQ, WF²Q, Lottery and Stride, BVT, VTRR

- **Two approaches**
 - **Global scheduling**
 - The system has a single global process queue
 - Processes are dispatched to any available processors
 - **Partitioned scheduling**
 - Each processor has a separate process queue
 - Each queue is scheduled by an independent scheduler
 - Process migration may be allowed or not

Global vs. Partitioned Scheduling



Global Scheduling



Partitioned Scheduling

Global vs. Partitioned Scheduling

- **Global scheduling** migration issue, cache problem
 - It is generally believed that global scheduling can achieve better performance
 - However, it can be inefficient due to the contention at the single queue and increased cache misses

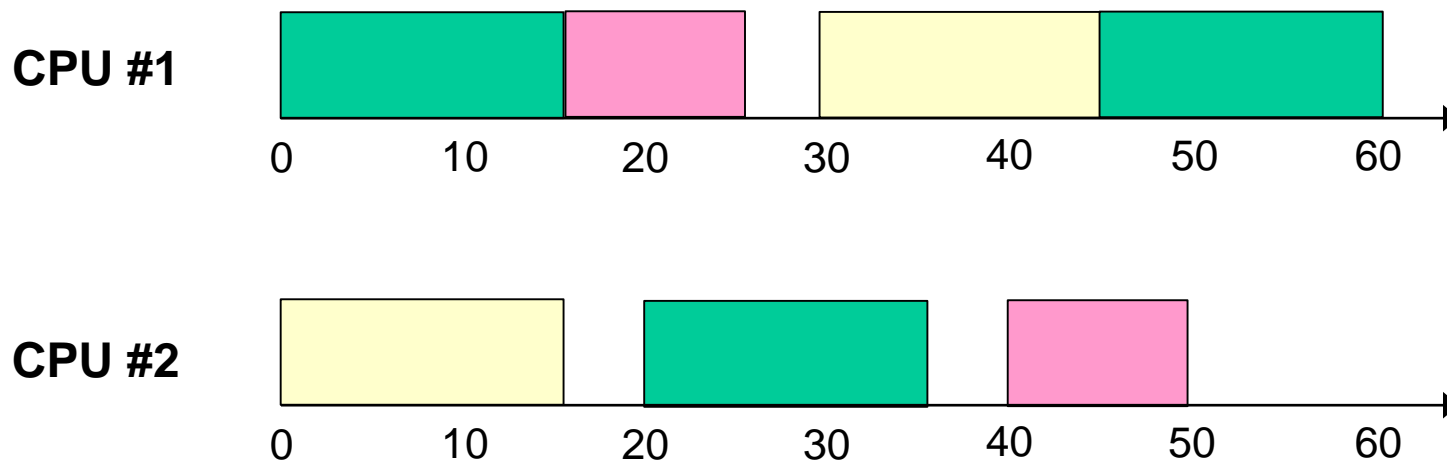
- **Partitioned scheduling** 실제로는 이 방식을 많이 사용한다. 그 이유는, RM scheduling 방식은 single core에서 증명된 스케줄링 방식이기 때문이다.
 - Performance can vary depending on the initial distribution of processes, i.e., a bin-packing problem
 - Different scheduling policies can be employed across processors
 - We can use the rich and extensive results from the uniprocessor scheduling theory

bin-packing problem => NP problem

Global EDF

➤ Consider the following tasks

- Process X ■: period = 20, WCET = 15, deadline = 20
- Process Y ■: period = 30, WCET = 15, deadline = 30
- Process Z ■: period = 40, WCET = 10, deadline = 40




Partitioned EDF


➤ Consider the following tasks

- Process X ■: period = 20, WCET = 15, deadline = 20
- Process Y ■: period = 30, WCET = 15, deadline = 30
- Process Z ■: period = 40, WCET = 10, deadline = 40

**CPU #1
(X)**



**CPU #2
(Y, Z)**



Schedulability Analysis

➤ Global EDF

- There is no single efficient test
- Most tests are very complex

Theorem 2 (GFB). A task set τ is schedulable with global EDF if

$$\lambda_{\text{tot}} \leq m(1 - \lambda_{\text{max}}) + \lambda_{\text{max}}. \quad (1)$$

Theorem 3 (BAK, from [3]). A task set τ is schedulable with global EDF if, for all $\tau_k \in \tau$, there is a $\lambda \in \{\lambda_k\} \cup \{U_\ell | U_\ell \geq \lambda_k, \ell < k\}$ such that

$$\sum_{\tau_i \in \tau} \min(1, \beta_{i,k}(\lambda)) \leq m(1 - \lambda) + \lambda, \quad (2)$$

where

$$\beta_{i,k}(\lambda) = \begin{cases} U_i \left(1 + \frac{\max(0, T_i - D_i)}{D_k}\right) & \text{if } U_i \leq \lambda \\ U_i \left(1 + \frac{T_i}{D_k}\right) - \lambda \frac{D_i}{D_k} & \text{if } U_i > \lambda. \end{cases}$$

Theorem 6 (FFDBF from [10]). A task set τ is schedulable with global EDF if $\exists \sigma | \lambda_{\text{max}} \leq \sigma < \frac{m - U_{\text{tot}}}{m - 1} - \epsilon$ (with an arbitrarily small ϵ), such that $\forall t \geq 0$,

$$\text{ffdbf}(t, \sigma) \leq (m - (m - 1)\sigma)t \quad (10)$$

It can be proved that it is sufficient to check only those values of t in $\{kT_i + D_i | k \in \mathbb{N}\}_{i=1}^n$ that are smaller than²

$$\frac{\sum_{\tau_i \in \tau} C_i \left(1 - \frac{D_i}{T_i}\right)}{m - (m - 1)\sigma - U_{\text{tot}}}. \quad (11)$$

...





...

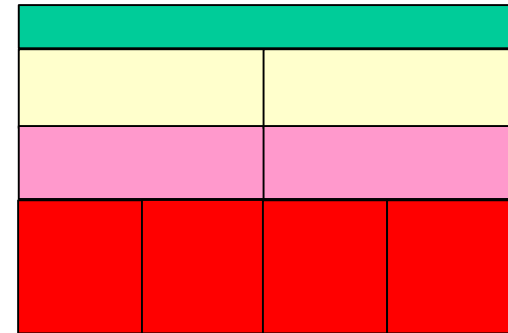
➤ Partitioned EDF

- Sufficient to check if the CPU utilization does not exceed 100% for each processor

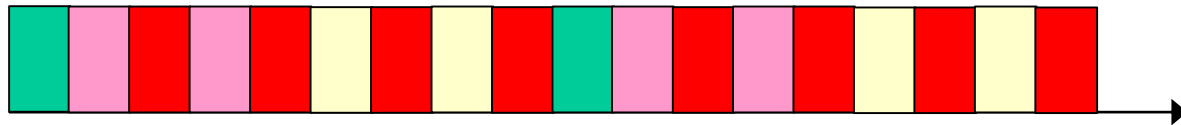
Global WFQ

➤ Consider the following tasks


- Process A : weight = 1
- Process B : weight = 2
- Process C : weight = 2
- Process D : weight = 4



CPU #1


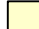





CPU #2



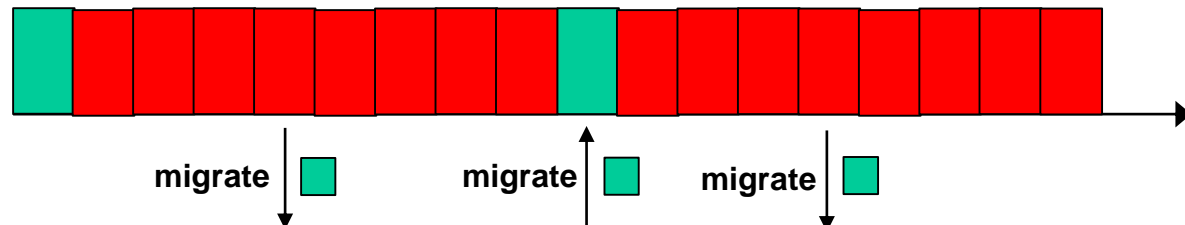
Partitioned WFQ with Load Balancing

➤ Consider the following tasks

- Process A : weight = 1
- Process B : weight = 2
- Process C : weight = 2
- Process D : weight = 4



CPU #1
(A, D)



CPU #2
(B, C)

